

Energy Berger

energyberger.github.io / CompSCI-630

Systems Lunch @ Fridays - Fall

Random number

TA Abhinav Jangda

Github!

Hot CRP

energy@cs.umass.edu

- Scribes!
  - course load :
    - 1 core / semester for MS
    - ..
    - 1<sup>st</sup> semester MS / PhD

> 1 core = too much!
  - schedule - yes, we have a midterm & final!
  - reviews
- 
- ## FORTRAN (in COBOL)
- 1954
- |  |  |
|--|--|
| military <ul style="list-style-type: none"> <li>- calculating</li> <li>- artillery trajectories</li> <li>- decryption</li> </ul> | first computers<br>— people!                   |
| general-purpose computing machine  | automatic computers                            |
| Turing machine<br>von Neumann machine  | ] data & code together<br>Harvard architecture |
|  | fetch-decode-execute cycle                     |

## assembly language

- lowlevel
- direct access to memory
- operations not abstract

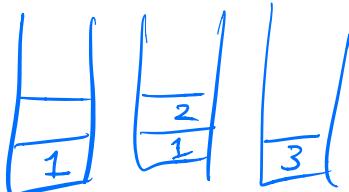
- stack architecture
- register architecture

A, B/C/D  
AX BX CX DX

- portability

ASM ≠  
~~portable~~

MOV 1, A  
MOU 2, B  
ADD A, B, C



PUSH 1  
PUSH 2  
ADD  
POP

"Intel" "ARM" (subset of instr.)  
JVM, .NET, Python bytecodes,

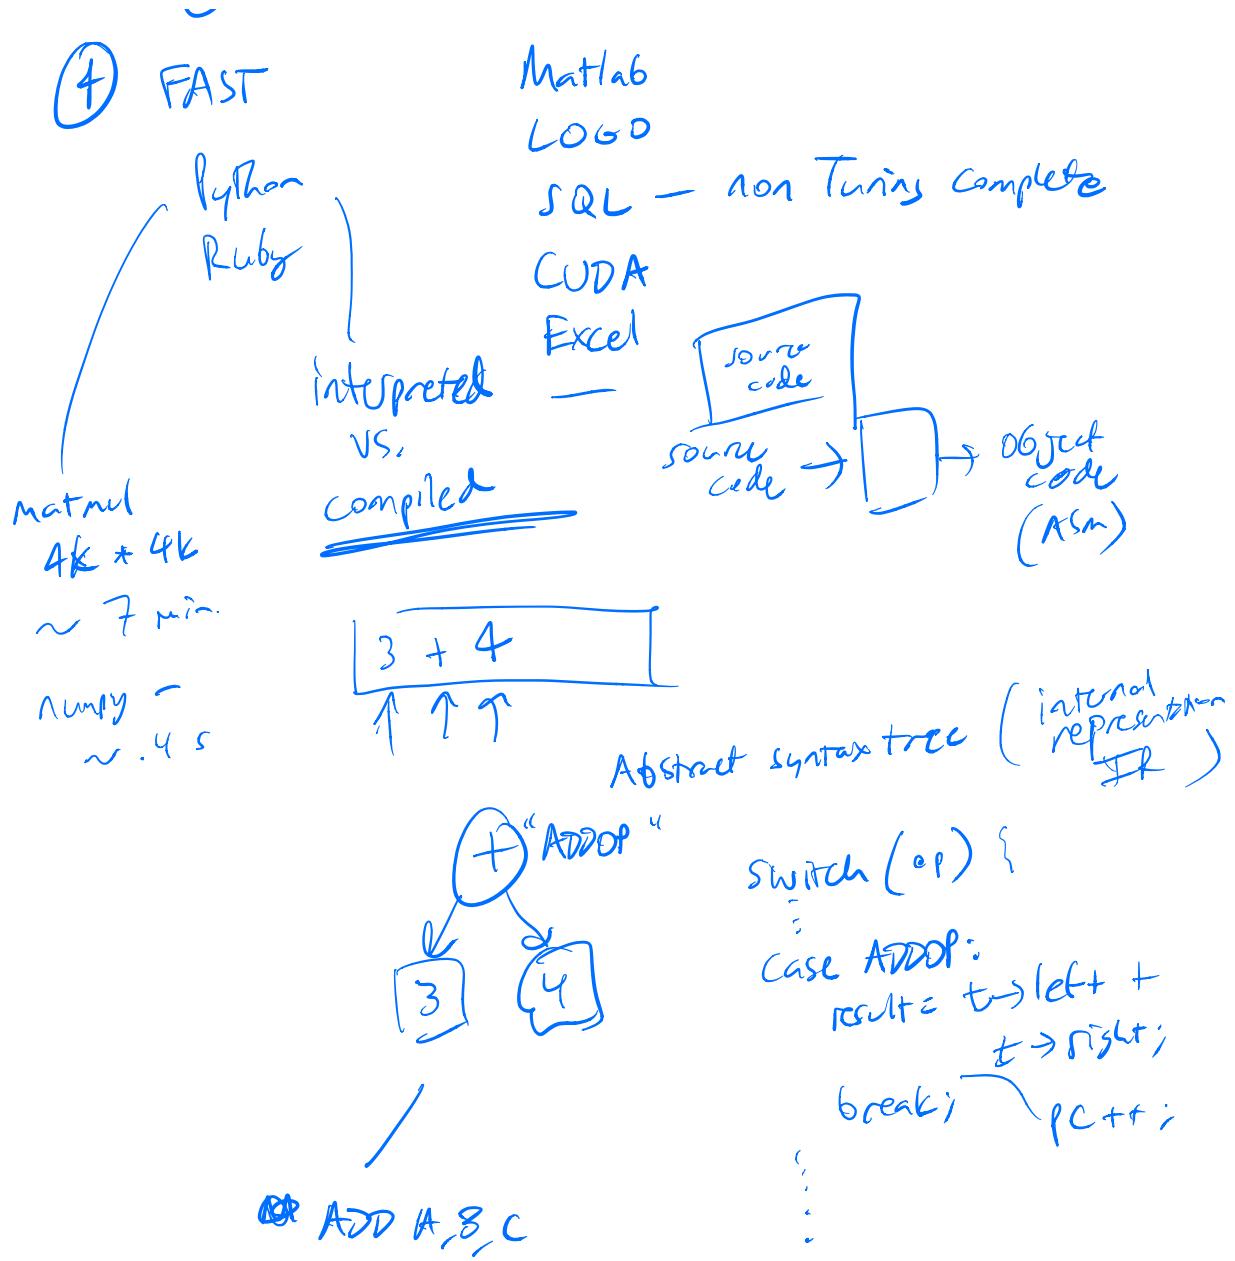
WASM

backward compatibility

- hard to read, hard to debug

- ① PORTABLE
- ② READABLE
- ③ WRITABLE

- COBOL - managers  
FORTRAN - scientists } domain-specific languages



programmer time MORE USEFUL than cycles

— modern POU

- productivity
- debugging
- ~ 100x faster!

interpreter → just-in-time compiler (JIT)

Python      PyPy  
JavaScript    V8  
                Chrome

WASM

Predicates -  
- Parsing  
- Compiling  
- program analysis  
- type coercion  
- scope  
- object orientation  
- modularity  
- structured prog.

FORTRAN  
1st optimizing compiler

I, J, K      integer  
X, Y, Z      float

DO 10 I = 10 10      \*

(was comma) line number

10 CONTINUE

BASIC



$\text{DO } \cup^1 \phi \cup^2 \psi = \cup^1 \phi$

control flow

$\text{DO } \phi \text{ I} = 1 \cup^1 \phi$

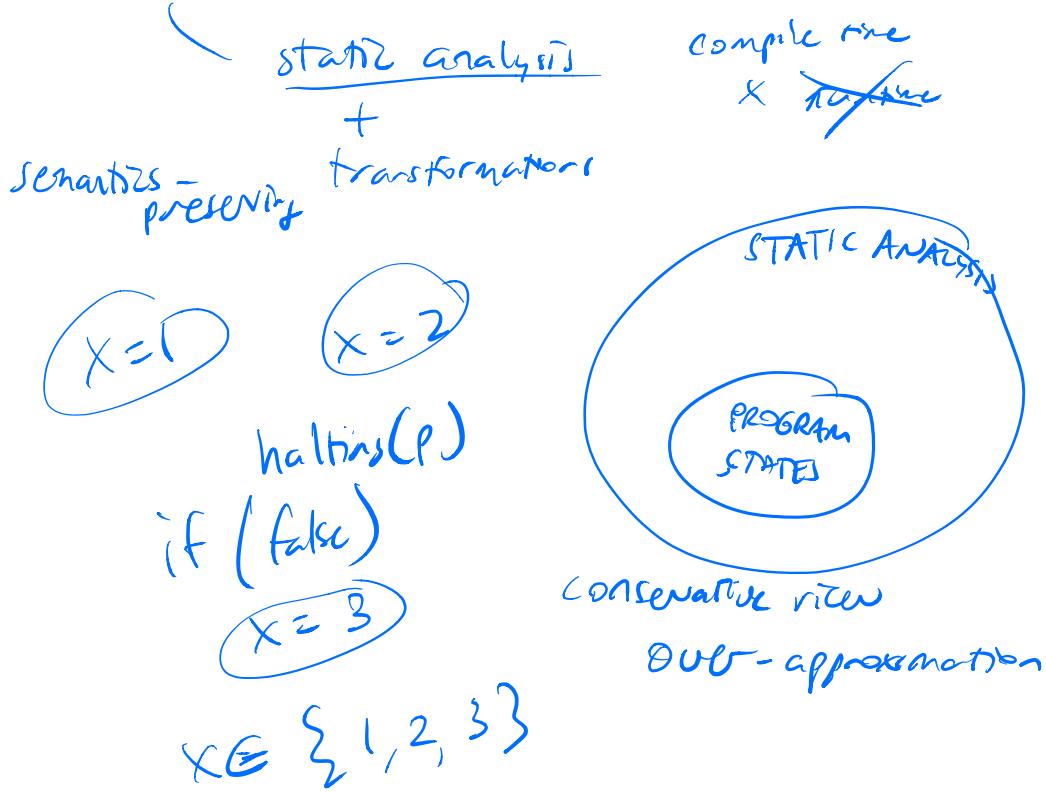
FORTRAN 77, FORTRAN 90 ~

writing

LEGACY  
FAST

The Fastest PC today! 2019 !!

optimization



## Halting problem

halts (program-text)  
→ true if program terminates  
                          <sup>always</sup>  
→ false otherwise

$p'$       if    halts( $p$ )  
                  run forever  
else            terminate

### ① Constant folding

$$x = 3 + 4$$
$$\Downarrow$$
$$x = 7 \quad * PI * PI$$

### ② Constant propagation

$$x = 7 \quad z = 12$$
$$y := x + 3 + z$$
$$\Downarrow$$

$$x = 7 \quad z = 12$$
$$y = 22$$

### ③ ``strength' reduction

$$x = x \times x$$

↓

$$x = x * x$$

/ %

expensive  
cheap

$$x = x \% 8$$

↓

$$x = x \& 7$$

\_\_\_\_\_ 111

### ④ vectorization

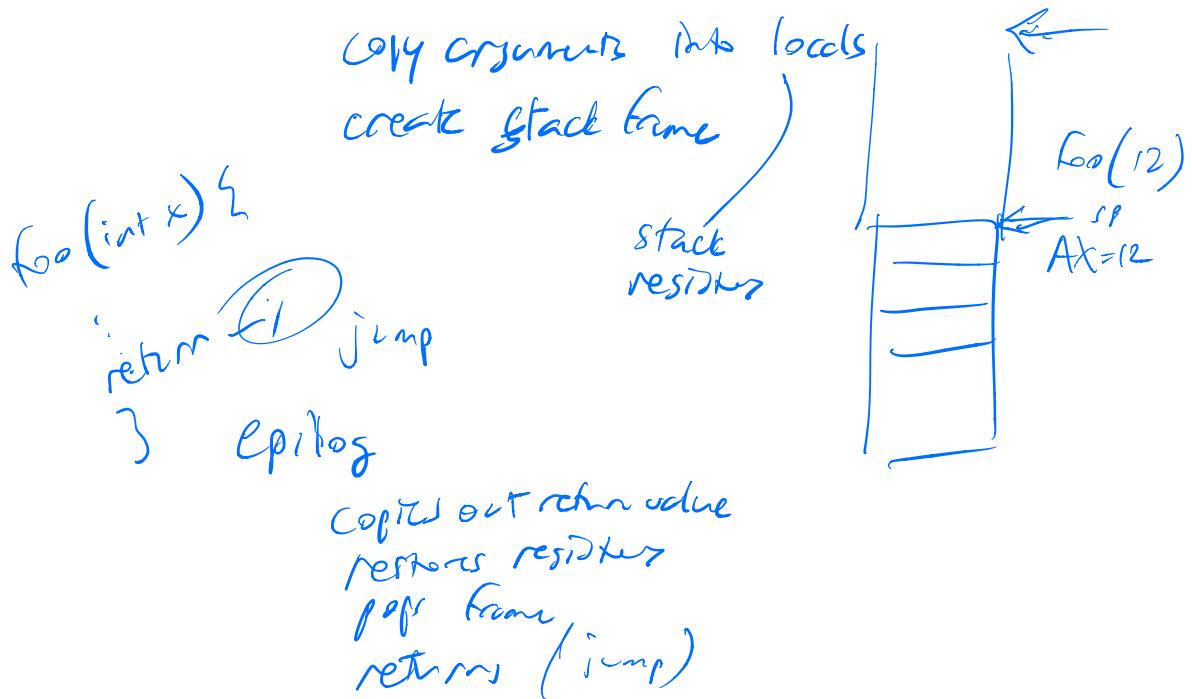
8	12	13	45
↓			
16	24	26	90

```
for (i=0; i<16; i++) {
    a[i] *= 2;
}
|||  

vector ops (4x speedup)
```

### ⑤ inlining

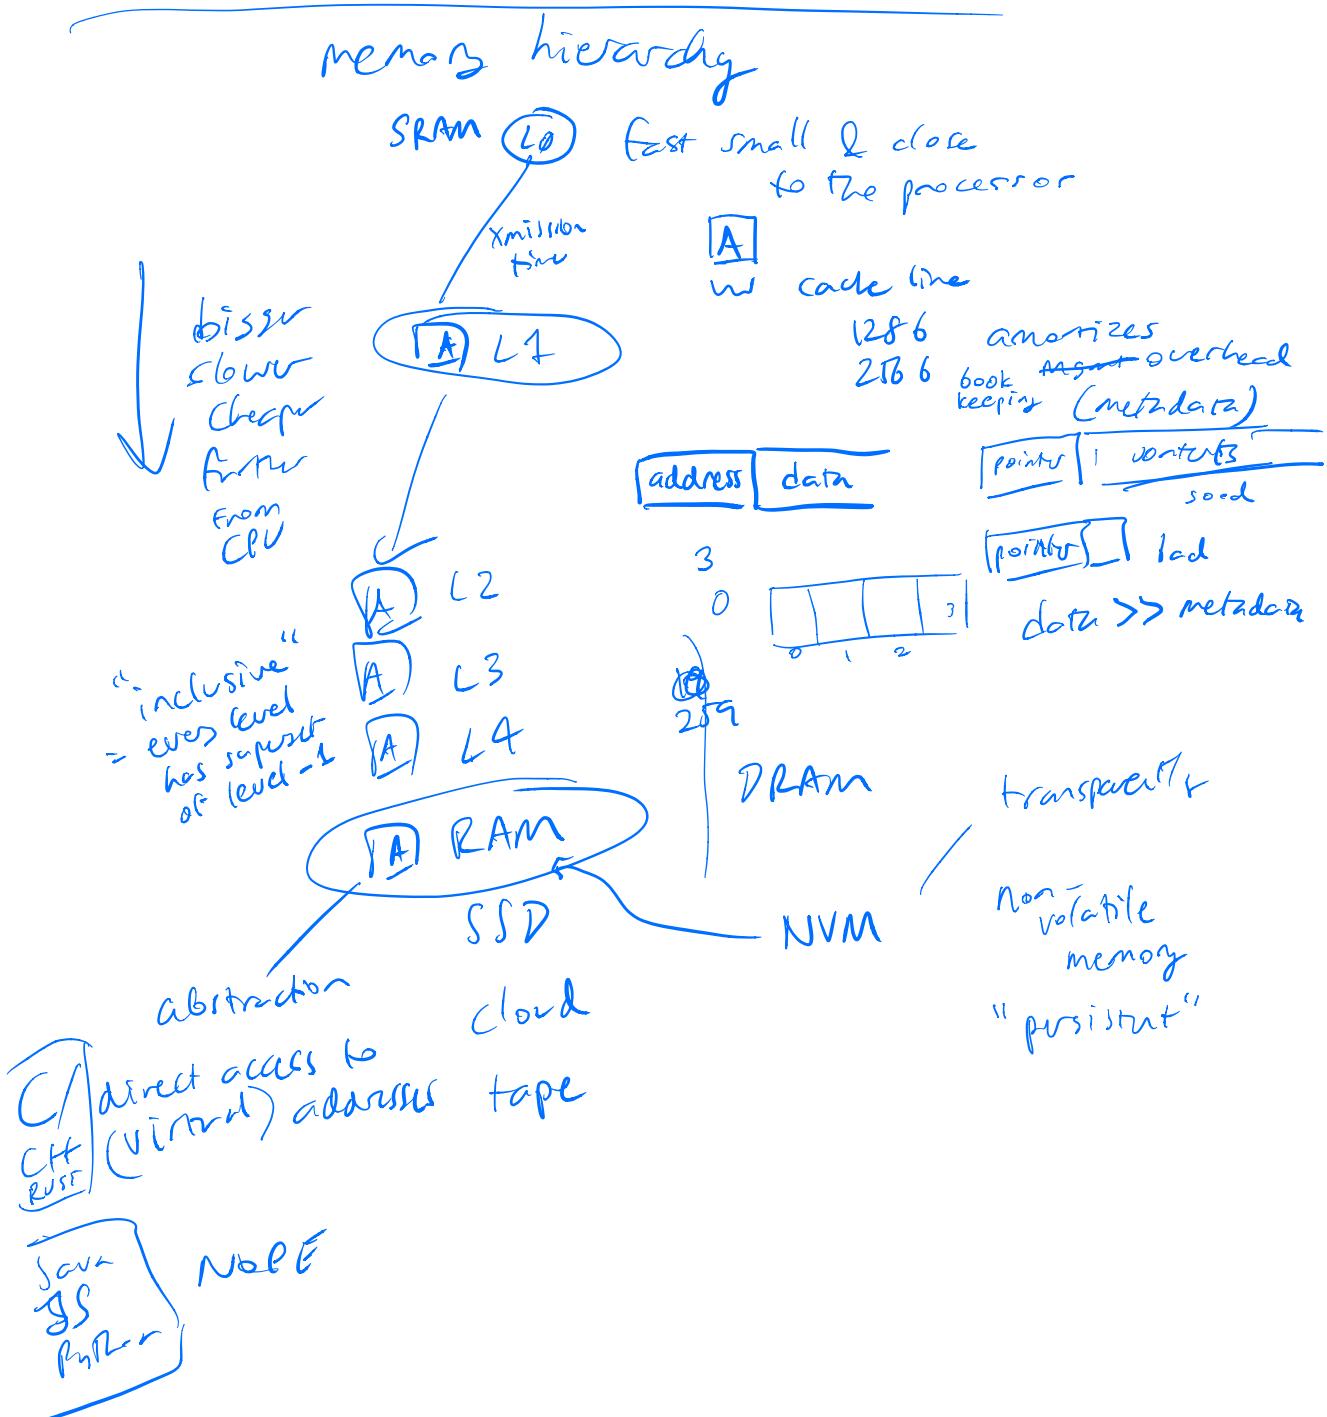
## function prolog

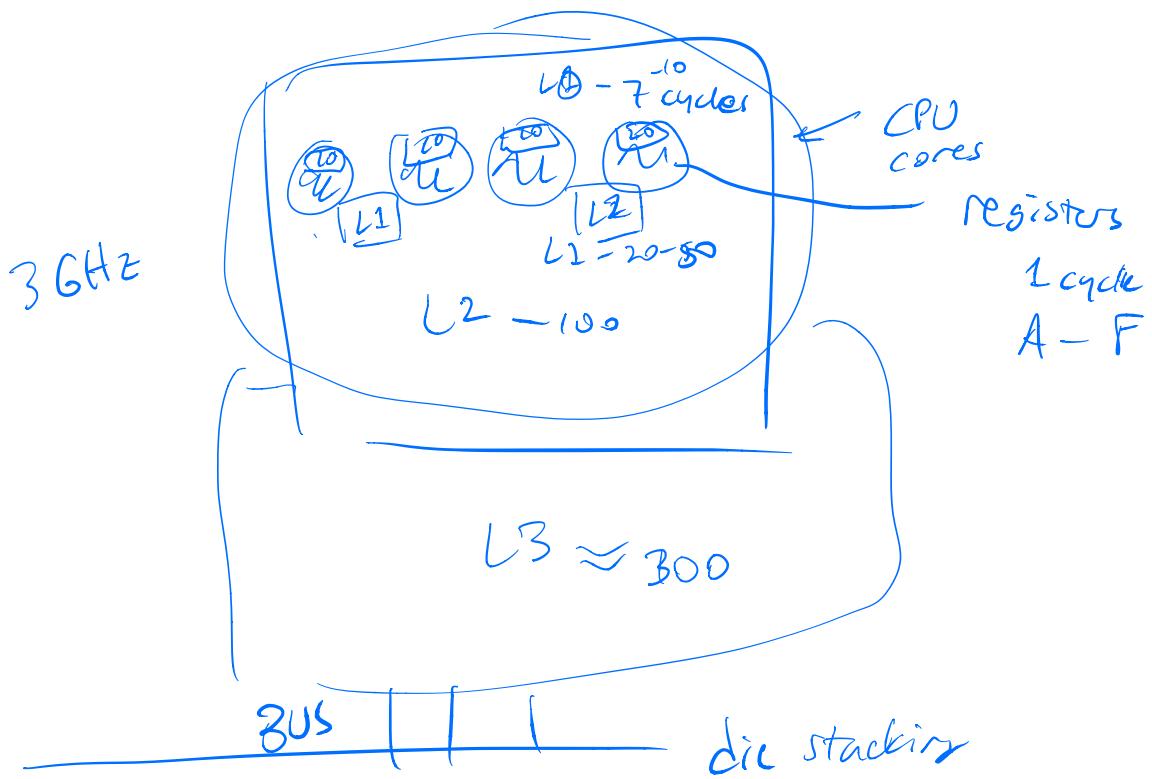


```
int add(int a, int b) {  
    return a + b;  
}
```

foo() {  
 z = add(a, b); ↴  
}  
↓  
foo() { z = a + b; }

- code size
- I\$ instruction cache





RAM "1 Gbs"

~~1000~~

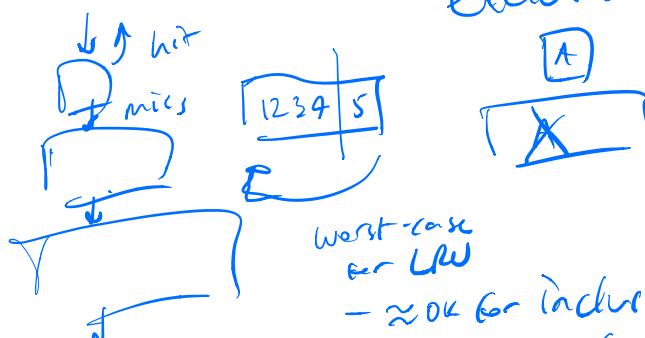
RAM bandwidth

1000  
10000

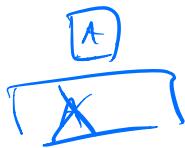
Caches — "replacement policy"

eviction "LFU"

latency  
wait time  
caches wide latency

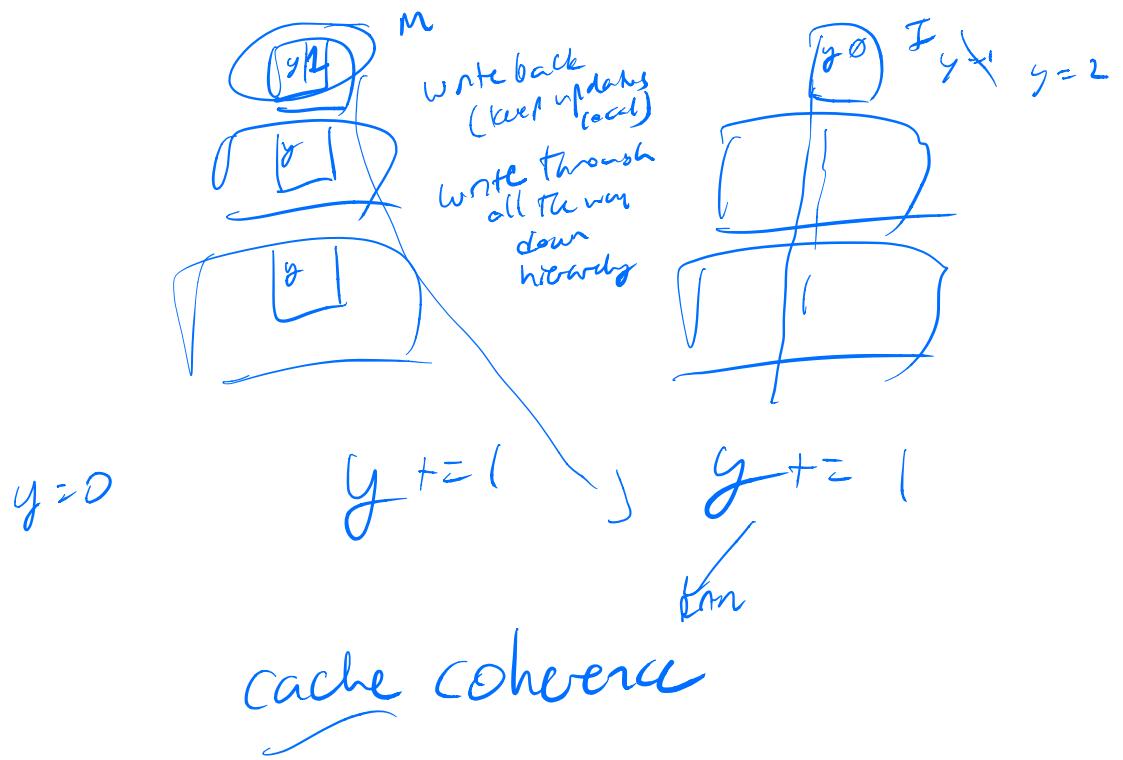


exclusive



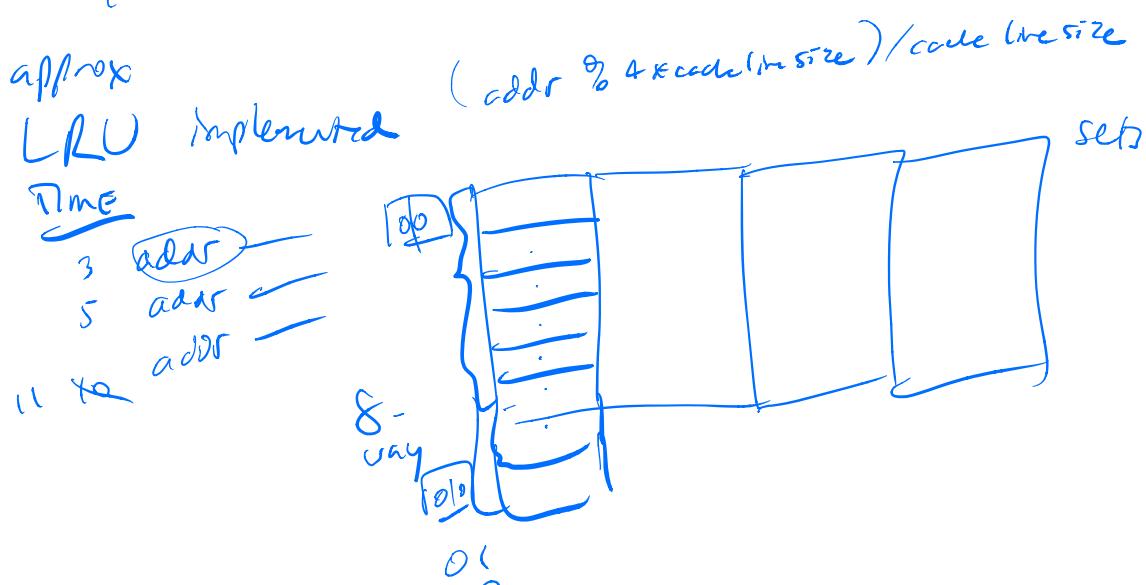
worst-case  
for LFU

- ≈ OK for inclusive
- catastrophic for exclusive



## MESI protocol

- { Modified
  - { Exclusive
  - { Shared
  - { Invalid
- I changed it  
 just me  
 several CPUs have it (modification)  
 trash



Cache misses

Can be different  
types -

Capacity - "hash collision"  
Conflict -  
Coherence - 3C's model

direct mapped - 1-way

2-way

4-way

8-way

16-way

fully associative -  $\infty$  way

TLB

translation  
lookaside  
buffer

Cache of  
virtual  $\rightarrow$  physical addresses

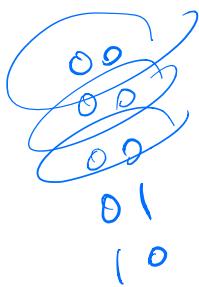
#include < stdio.h >

```
int main() {
    int a;
    printf("a = %p\n", a); 0x10000000
    return 0;
}
```

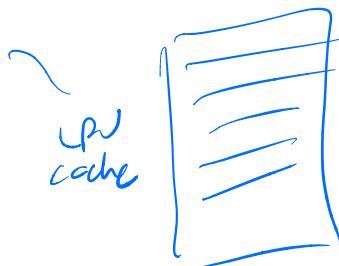
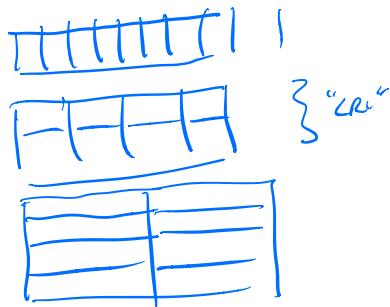
address

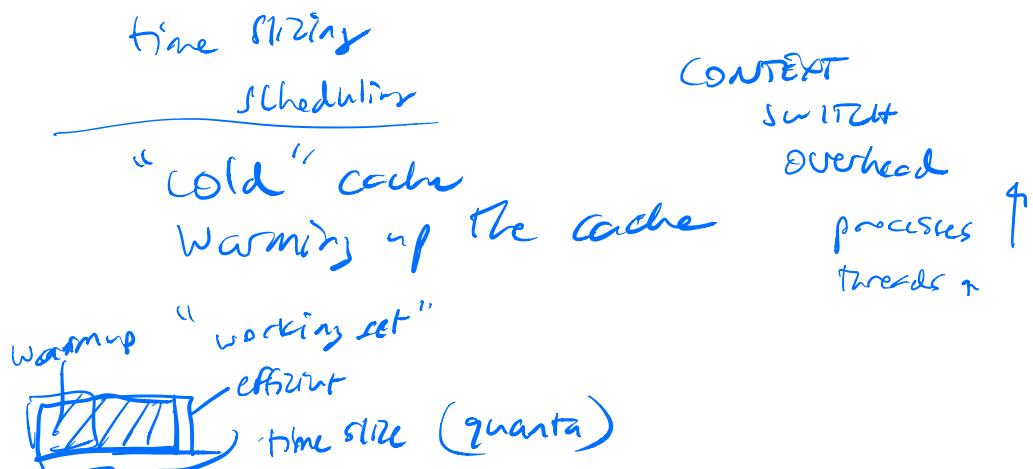
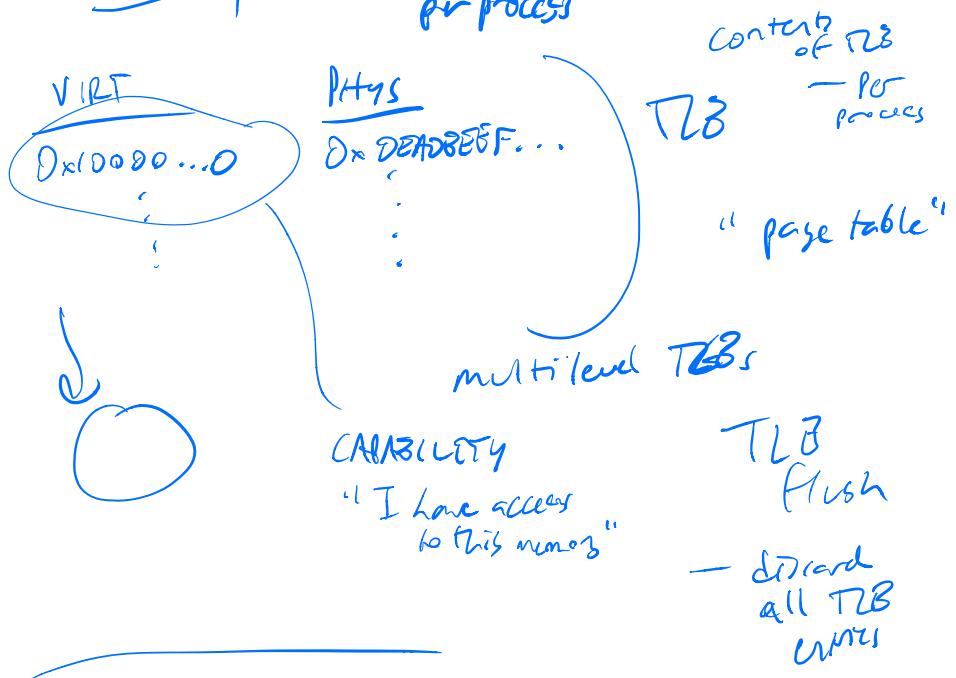
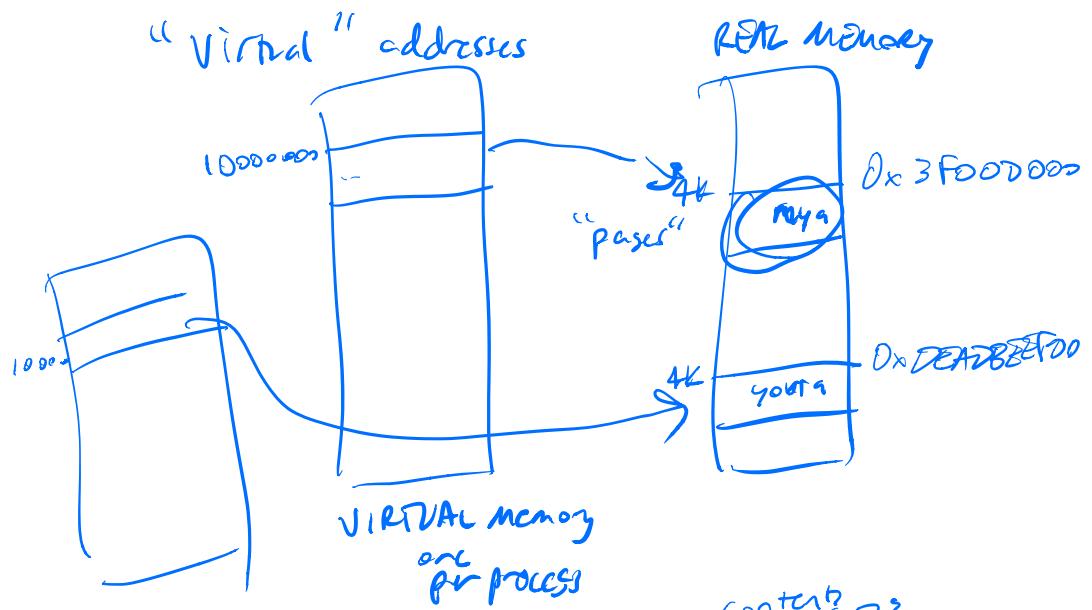
process isolation

10  
11



(11)  
(11)  
saturated  
counter







anomalous cleanup  
across the quantum



REGS < QUANTUM

register allocation

"spill" - copy  
contents  
of register  
to mem "eviction"!

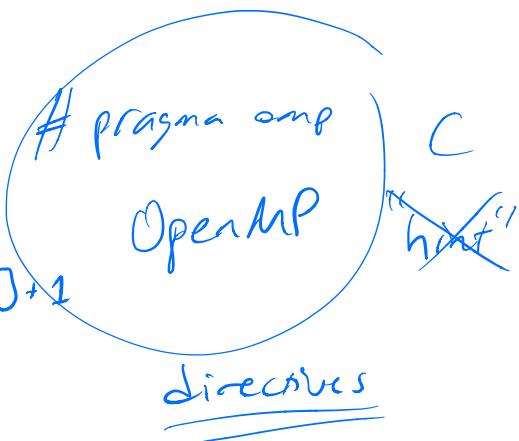
straight-line code  
- predictable!

"graph coloring" NP complete!

Why is FORTRAN fast?

~~OMP PARALLEL~~

```
DO I = 1, 100
  DO J = 1, 100
    A[I][J] = A[I][J-1] + 1
  ENDDO
ENDDO
```



automatic  
parallelization

"holy grail"

main thread  
 S  
 ↓  
 for ( $i=0; i < N; i++$ ) {  
     foo();  
 }  
 ↗  
 are there  
 loop carried  
 dependencies?

P  
 }  
 for ( $i=0; i < N; i++$ ) {  
     t[i] new thread (foo);  
 }  
 ↓  
 foo<sub>0</sub>, foo<sub>1</sub>, ...  
 [t[i].join()]

$\text{result}(S) \equiv \text{result}(P)$

parallel prog  
semantically equivalent  
to seq prog.

does `foo()` have side effects?

foo()  
 {  
     x++  
     (<sub>global</sub>)
 }

read  $x$  into a register  
 increment  $x$   
 write the result back

$x = 0$

A  
 R<sub>0</sub> ← x  
 R<sub>0</sub> ← R<sub>0</sub> + 1  
 x ← R<sub>0</sub>

$A ; B \Rightarrow 2$

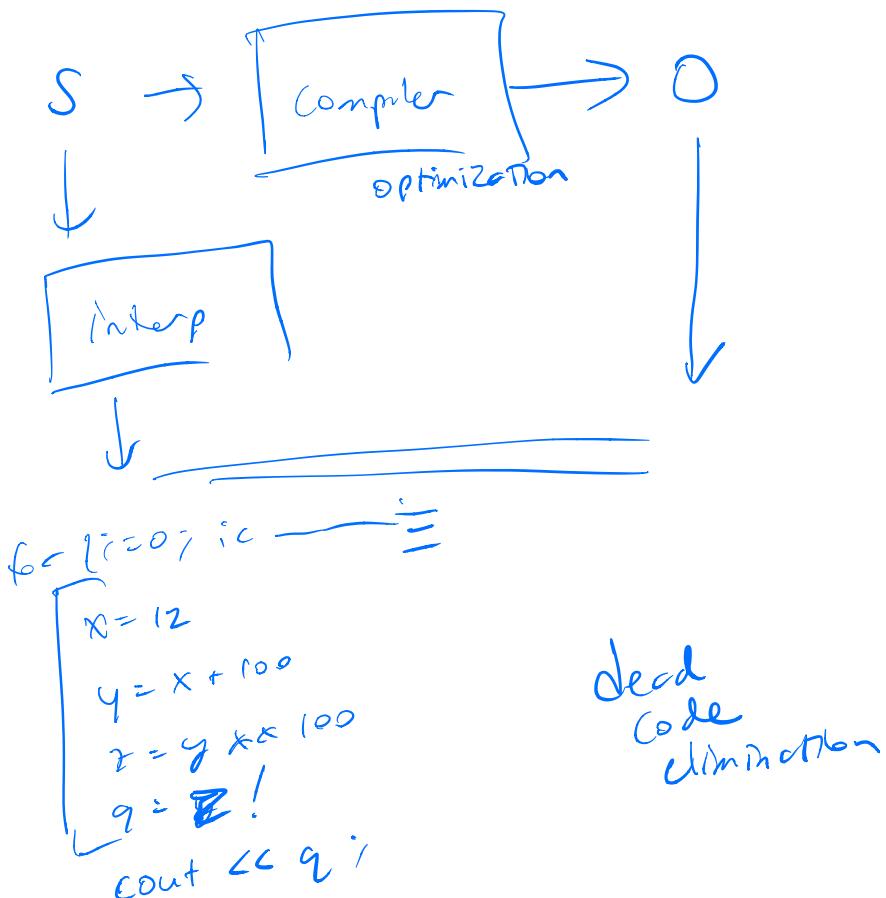
B  
 R<sub>0</sub> ← x  
 R<sub>0</sub> ← R<sub>0</sub> + 1  
 x ← R<sub>0</sub>

$$X \in \{1, 2\}$$

non-deterministic!

Race condition

Concurrency errors



```

foo()
{
    x =
}
  
```

```

foo (int i)
{
    i
    return ack(i, i+1);
}
  
```

Ackermann

```
foo (int *int i a,) {  
    a[i] = ack(i, i+1);  
}
```

```
foo (int * a, int *int i b) {  
    a[i] = -;  
    b[i] = a[i] + 2;  
}
```

```
int x[200]; int *y = &x[0];  
v = foo(x, y, i)
```

### aliasing

$p \{ \quad \}$       alias analysis  
 $q \{ \quad \}$       pointer analysis

$p \cap q = \emptyset$        $\text{int } *p; \quad p \in \{ \}$

if ( $\_\_$ ) {  
  $p = q;$        $p \in \{q\}$

} else {  
  $p = r;$        $p \in \{r\}$

symbolic execution

abstract interpretation  $\rightarrow$   
 $\wp \in \{q, r\}$

$$\begin{aligned} x &= 1^2 \\ x &= x - 100 \\ x &= \dots \end{aligned}$$

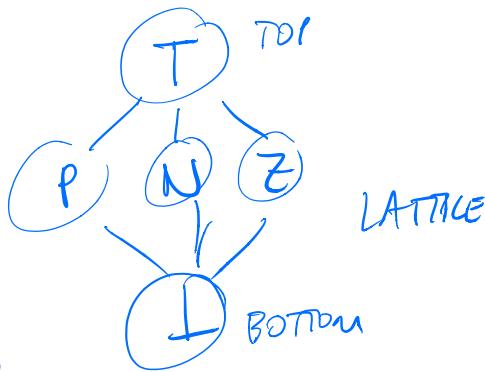
if ( $x > 0$ ) {  
 |  
 }

$x \in O, A, P$

$x = \_ \quad \{O, N, P\}$   
 $x = x - 1 \quad \{\}$   
 does this run over?  
 $O \Rightarrow O$   
 $N \Rightarrow P$   
 $P \Rightarrow N$

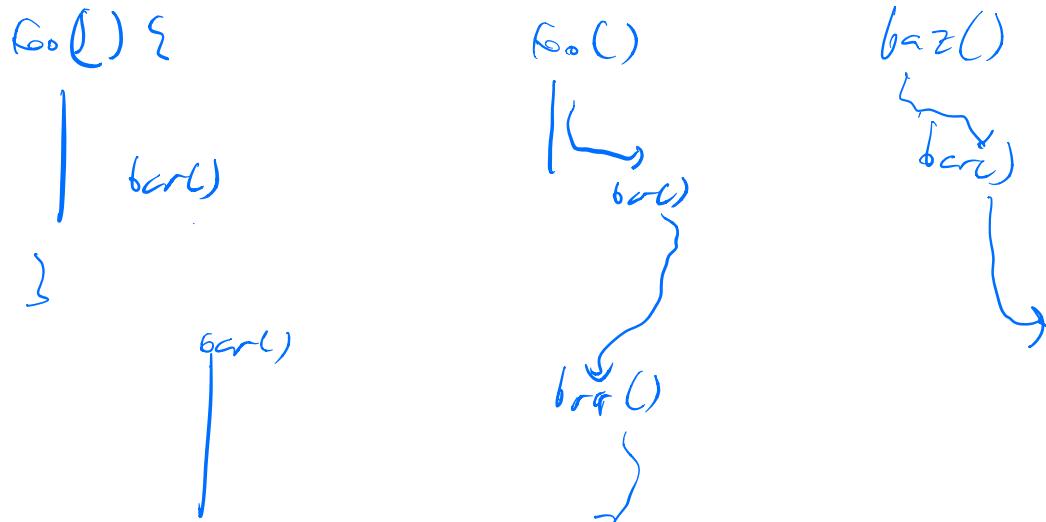
CONSERVATIVE APPROX  
 PROG BEHAVIOR

$\{x : T\}$   
 if ( $\neg$ ) {  
 $x = -1$   
 else  
 $x = 1$   
 }  
 $\{x : \perp\}$



$$\begin{array}{c} \text{if } x : P \\ \text{else } x : P \\ \hline x : P \end{array}$$

## intraprocedural vs. inter procedural



flow sensitivity

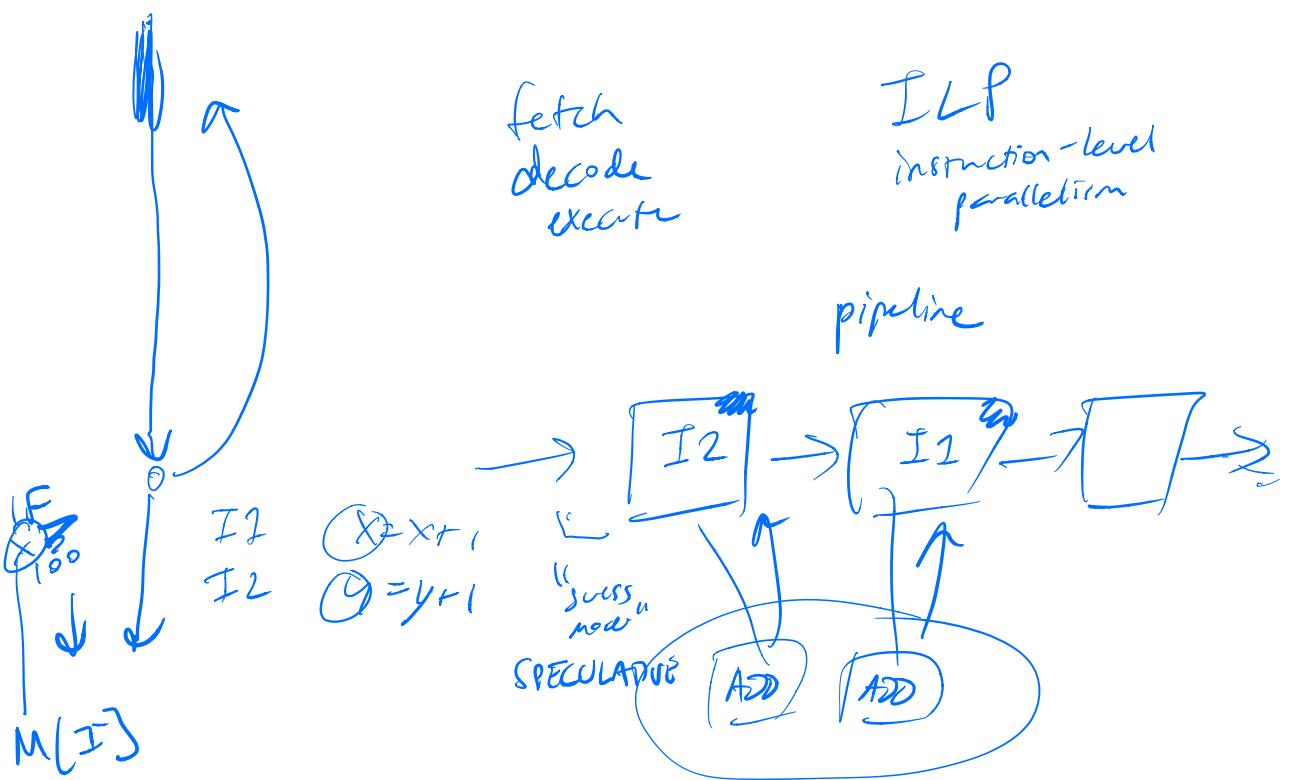


Context sensitivity

Steensgaard's algorithm

PCR sensitivity

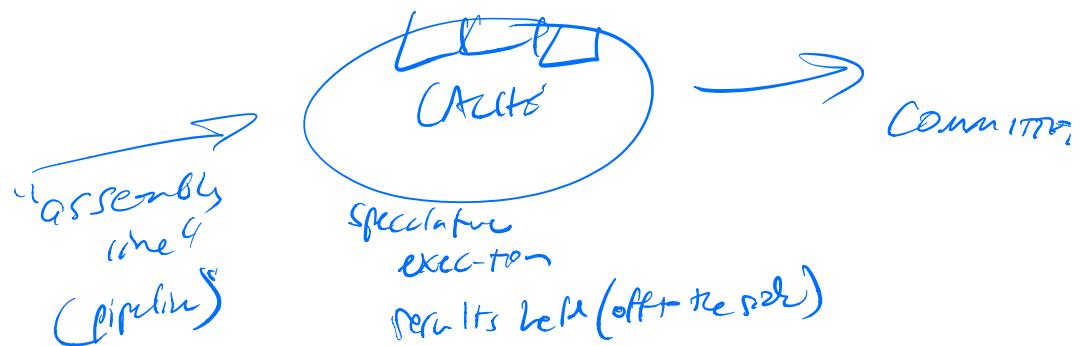
false positives false negatives $\neg W$	precision - recall tradeoff $V \sim \{\emptyset\}$ $\text{if } (x == \underline{\quad})$ $y = x / @ V$
--	--



Pipeline stall

"bubble"

"BACK EDGE DATA"



Speculative execution  
hide memory latency

# SPECTRE / Meltdown

timing channel

Covert channel

side channels

last branch taken

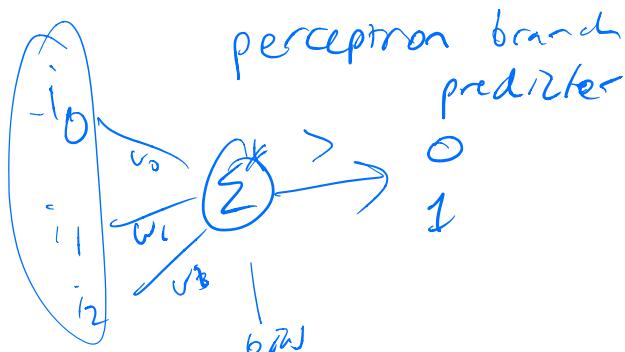
85-90% effective

PHAT

Modern

branch predictors

99.9% effective



---

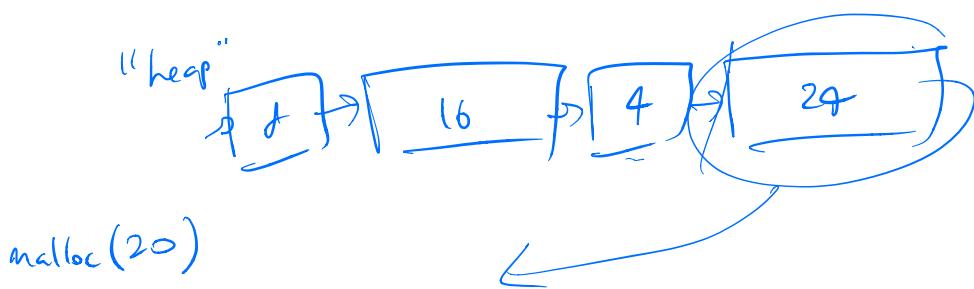
dynamic memory allocation

globally

stack variables

dynamic memory

"the  
heap"



Object 24

First-fit → SPACE INEFF

best-fit

SPACE  
EFFICIENT

slow ...

$$n \cdot O\left(\log \frac{M}{m}\right) * n$$

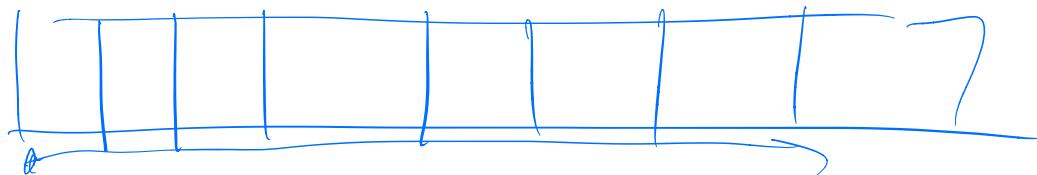
M - biggest  
m - smallest

$$O(M \times m)$$

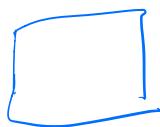
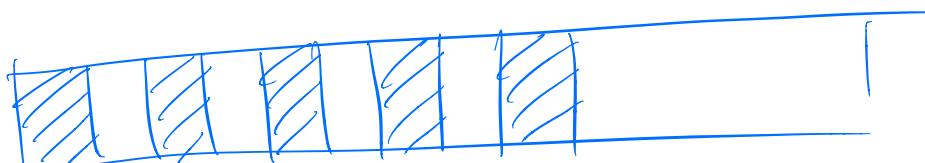
best fit

worst-case  
"fragmentation"

$$= \frac{\text{mem consumed}}{\text{mem requested}}$$



D → D → D



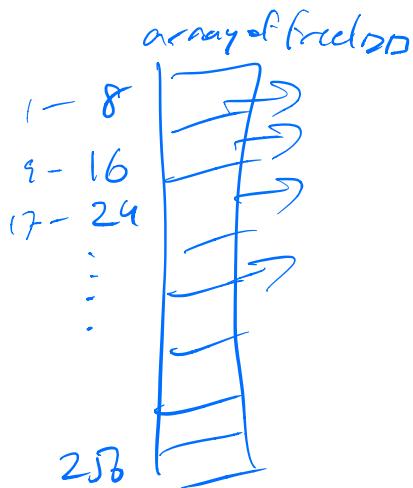
Compaction  
NOT IN C/C++

(L1E')

## Python, Ruby

"custom allocators"  
"ad hoc allocators"

malloc - "too slow"  
Reconsidering Custom Mem. Allocation



Py-object  
- malloc

if ( $sz \leq 28$ ) {

use small obj.

else

use malloc

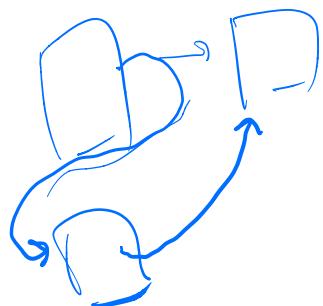
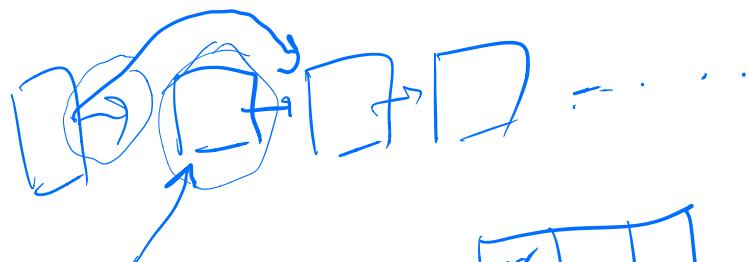
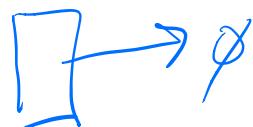
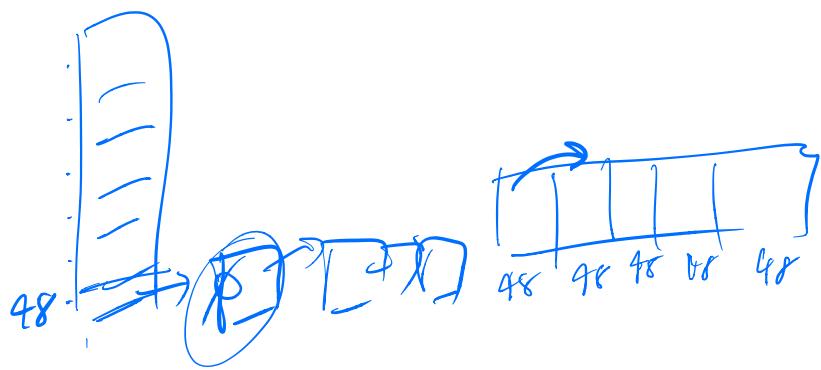
int  $\approx 32$  bits      C      Python  
long  $\approx 64$  bits      48      24/28

"a"        
1/16      50

{ "a": }  
~~2~~

240

60 20 75 100 140

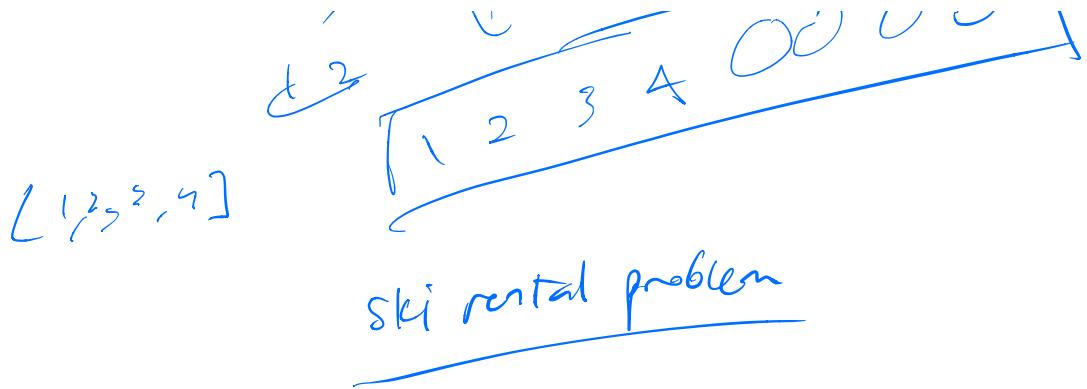


48			
ref count			

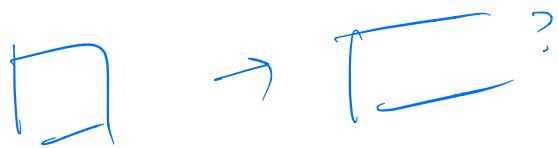
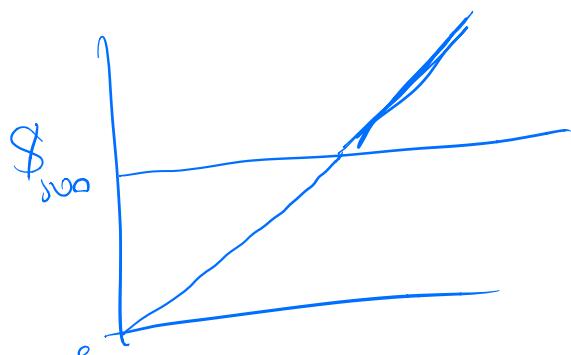
ref count

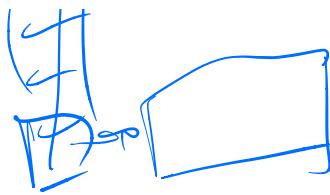
$$x = [1, 2, 3]$$





\$50 rent  
\$500 buy  
rent until  
you reach  
cost of buying

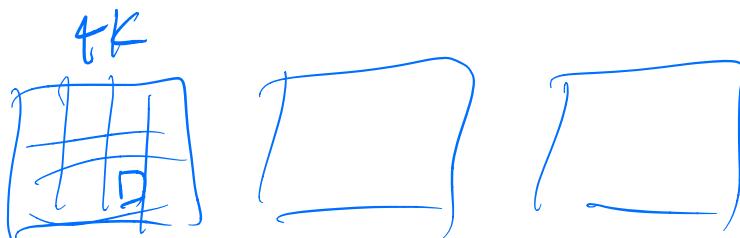
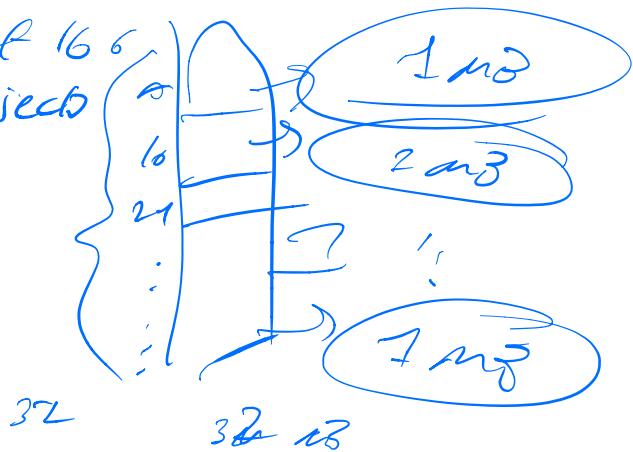




allocate 1 MB of 8 byte objects  
free them all

allocate 1 MB of 16 byte objects  
free them all

actual  
MMT 1 MB  
32x



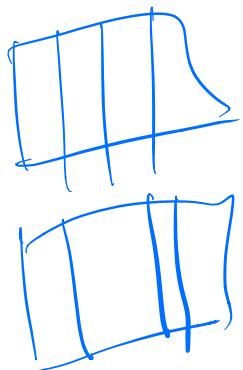
OS pages

Space-time tradeoff

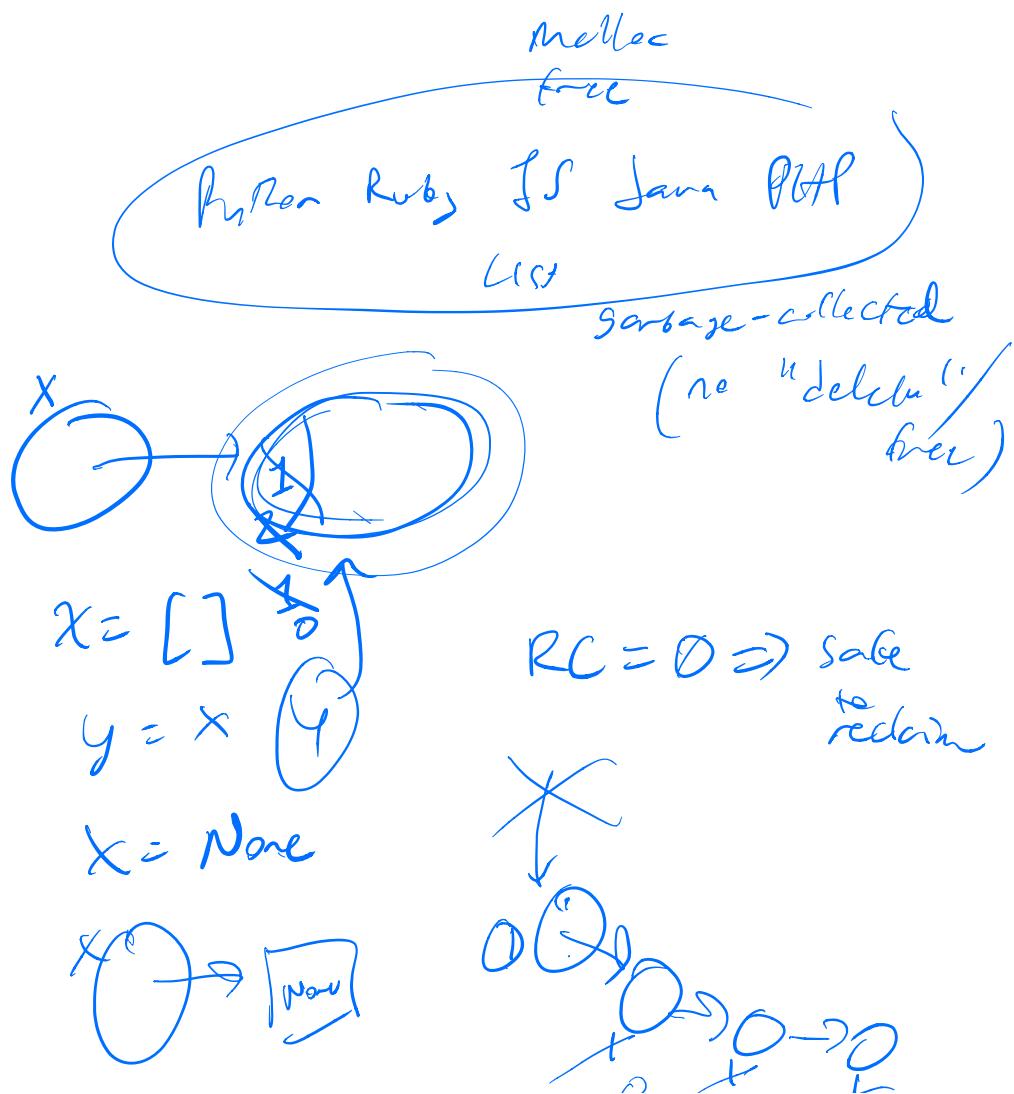
$$f(a) =$$

$$\begin{aligned} f(0) &= a \\ f(1) &= b \\ f(2) &= c \end{aligned}$$

0	a
1	b
2	c

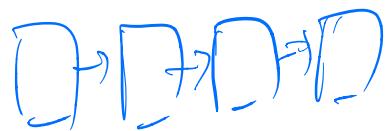


## Reference counting



$y = \text{None}$

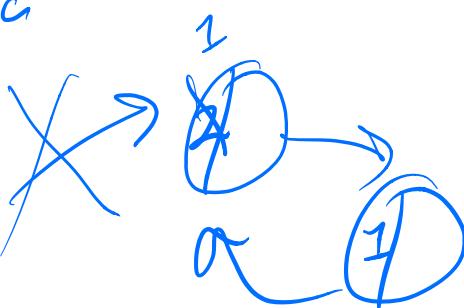
$\leftarrow \circ \rightarrow$



doubly-linked (2+)

$a = \{\text{'key'}: 0\}$  CYCLE

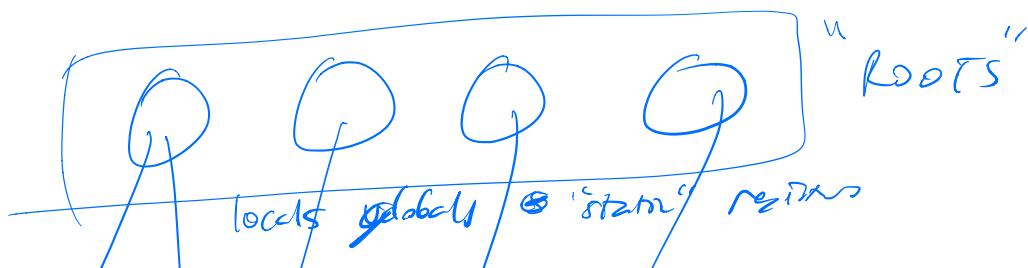
$a[\text{'key'}] = a$



HC  
Incomplete  
(1956)  
GC algorithm

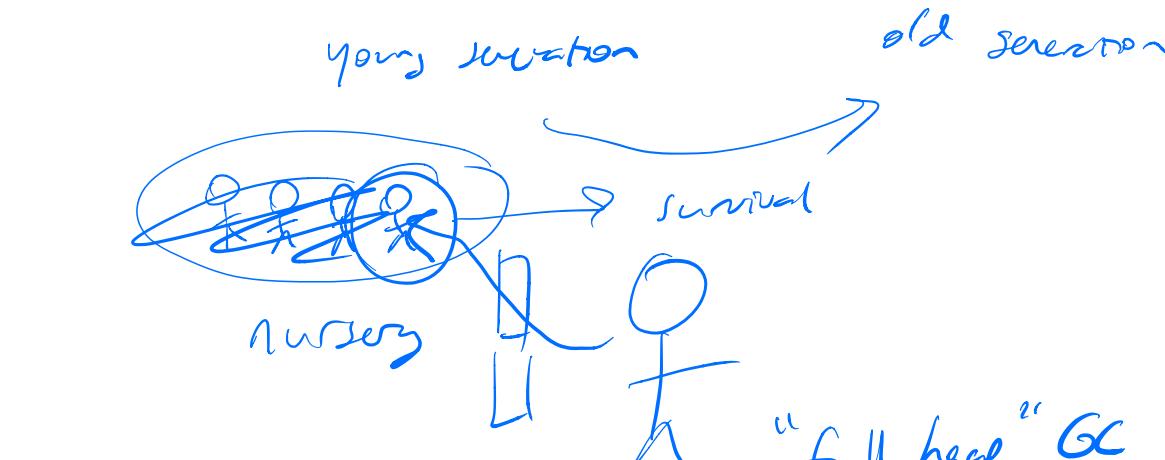
---

Mark-sweep (1978)  
"generational" GC (1983?)

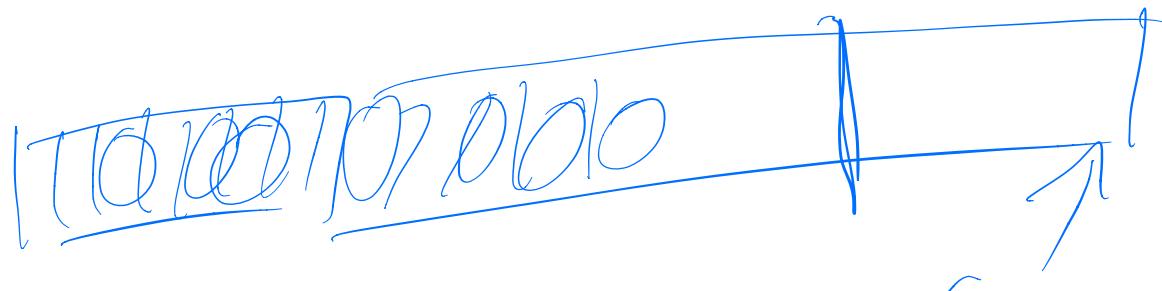
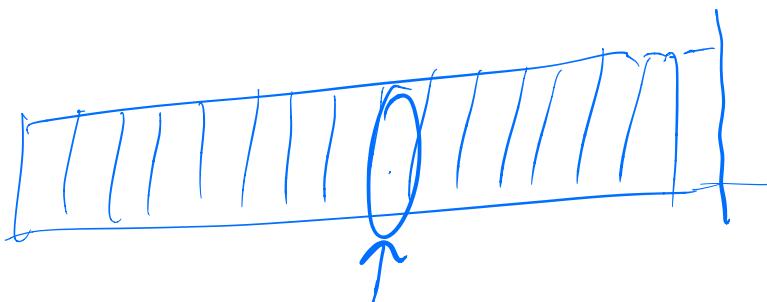




most objects die young



gen GC →  
takes pressure  
off the full  
heap GC

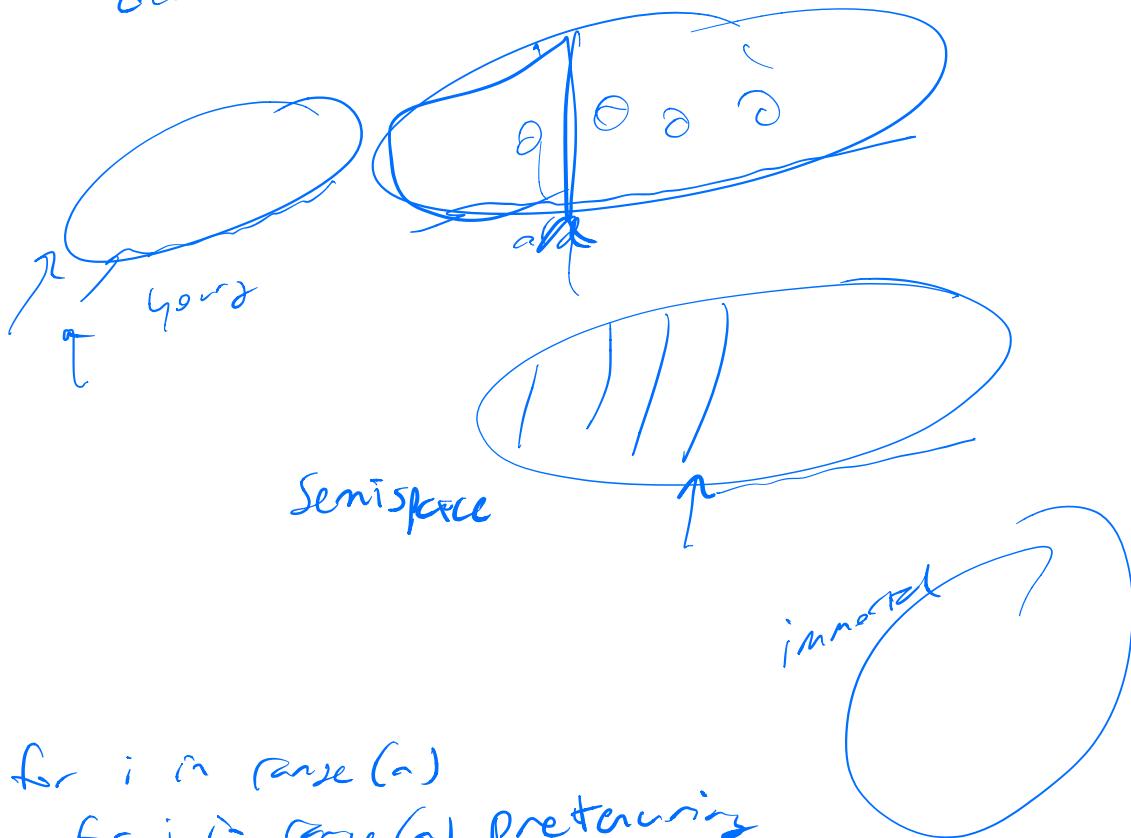


GC

$\sim 2 \rightarrow 5x$

"live  
objects"  
"dead"

Quantizing the cost of  
GC vs. Explicit deallocation



for i in range(a)

    for j in range(n) pretenuring

        for k in range(n)

$$C[i][j] += \delta A[i][k] + \delta[k][j]$$

• pyc

PyPy

# Ontogeny recapitulates phylogeny

SCRIBERS

Lillian  
& Sam

OS: single user, single process, no protection

multiple users, multiple processes

"time-sharing"

time slices

memory protection

file system protection

:

Mac 1984

1 user

1 process

no protection

VM

"MMU"

memory  
management  
unit

PC

Apple  
IBM PC

:= 1981

Motorola 6802  
Intel 8080...  
68000 Macs

) no MMU

MS-DOS

Microsoft

Microsoft BASIC

Altair

Gates & Allen

CP/M

8080

Z-80

8088

Gary Kildall

latencies

"security  
through  
obscenity"

AT&T

Unix

BSD

OS X

iPhone

stripped down Unix!

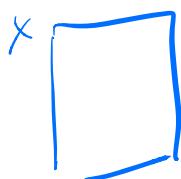
→ multiple processes  
mem protection  
security model!

PL?

→ dynamically typed lang — Values have types  
(statically typed lang) — Variables have types

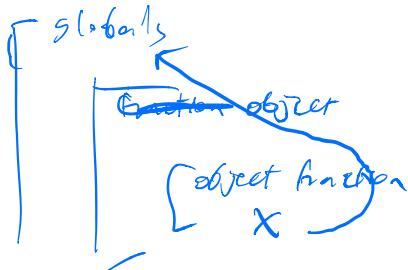
$x = []$   
 $x = {}$   
 $x = ""$

→ dynamically scoped



static / lexically scoped

"weird  
scope  
mics"



Match the innermost scope

foo() {

~x~

}  $x \leq x+1$

walking call stack

bar() {

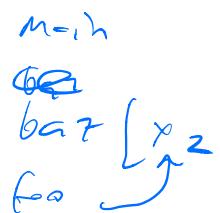
int ~~x<sub>1</sub>~~  
foo();

}

bar() {

int ~~x<sub>2</sub>~~  
foo();

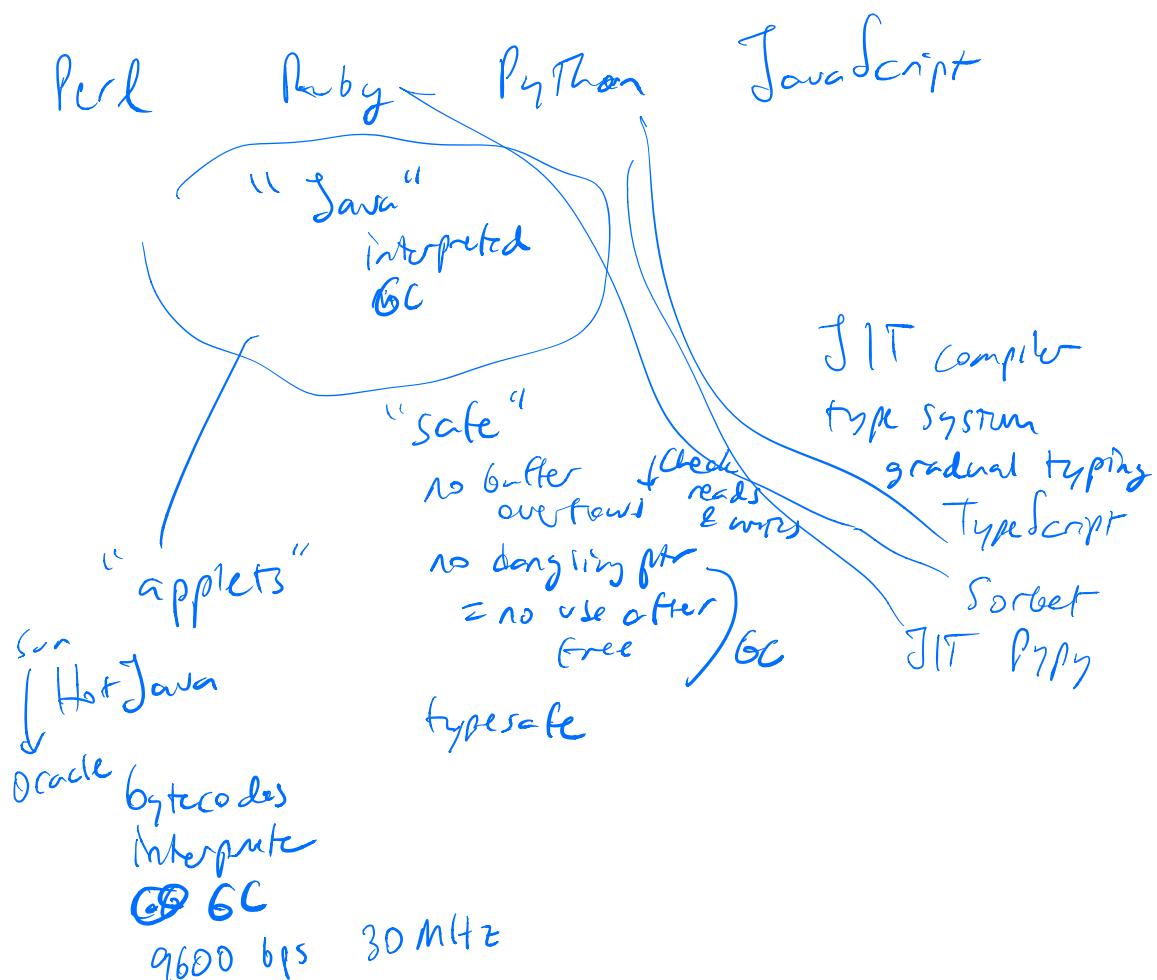
}



→ Interpreter eval

hard to compile  
into efficient code

— reference country GC



Netscape → Mozilla  
↓ Navigator

1995      Brendan Eich

Mozilla      {      }  
"curly-brace languages"

LISP Scheme  
( ( ( ) ) )

LiveScript      10 days

Browsix

Doppio

asm.js

subset that is easy to compile

asm.js  
 $x = (v \ 1 \ \emptyset);$   
 $x = a + b;$

Web Assembly

- safe
- efficient
- fast to parse!!



-- loading form compile connection

- ISA instruction set architecture
  - arbitrary memory (module mem protection)
  - arbitrary structure jmp GOTD
  - hard to analyze

Register-based  
stack-based  
architecture

A8  
CD

NaCl  
p NaCl

Compiler  
Smart  
"safe"  
subset  
of x86  
ARM

|||  
PUSH  
POP  
ADD

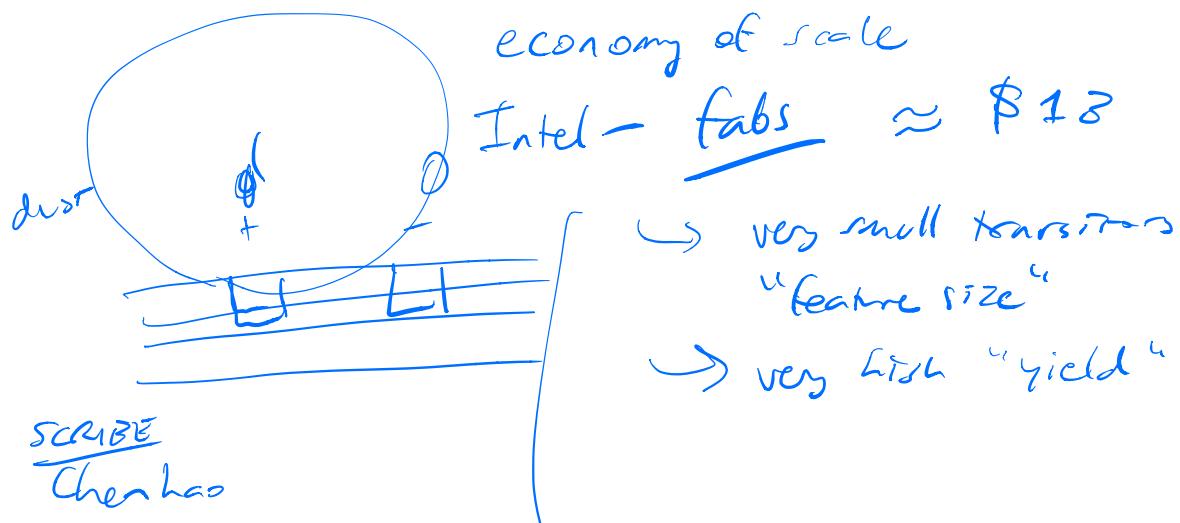
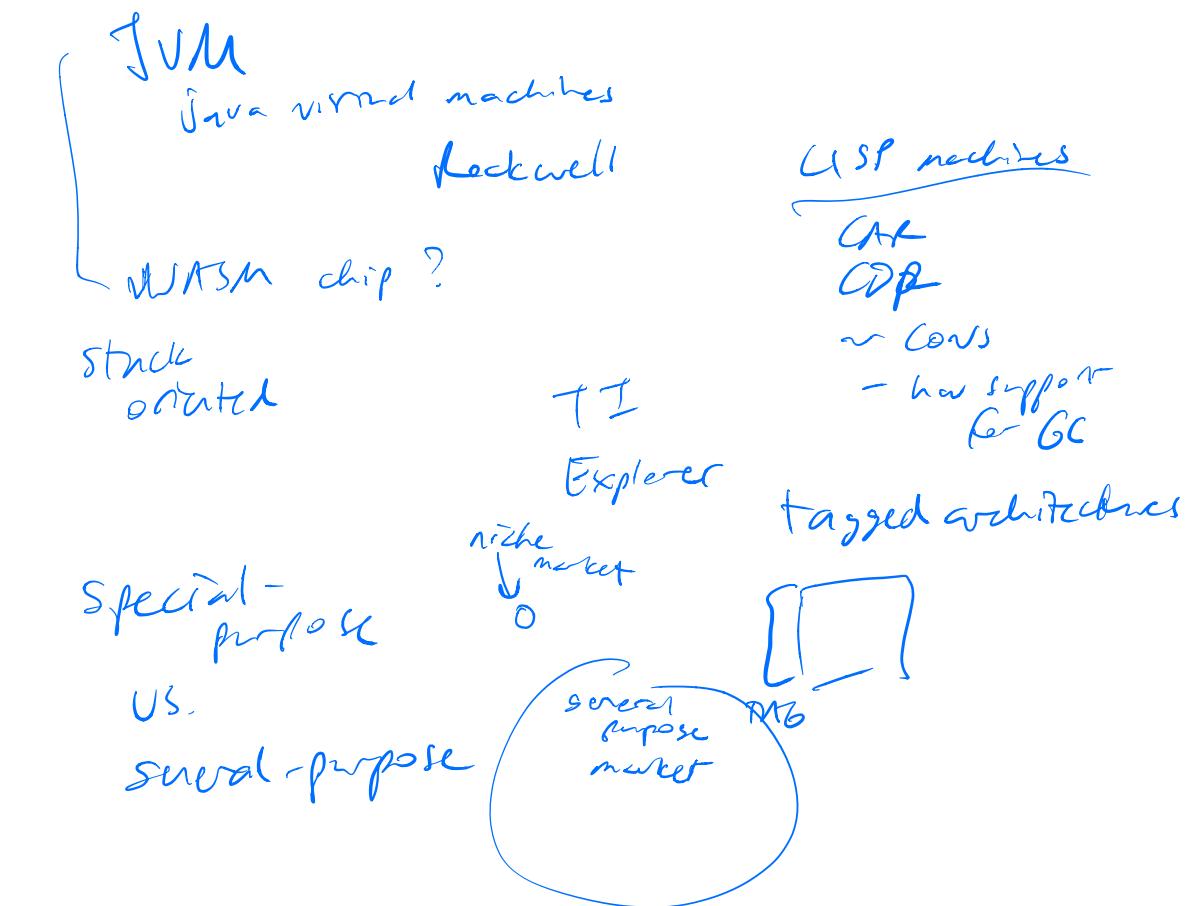
ROP  
Return-oriented  
programming  
Return-to-libc



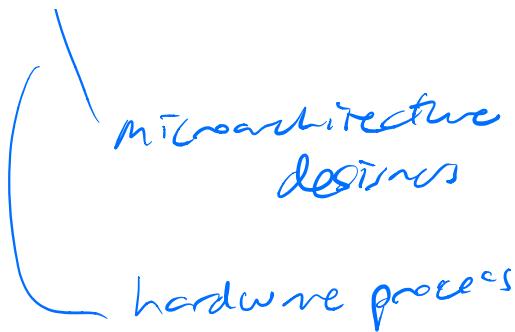
CFI  
Control-  
Flow  
Integrity

# WASM - CISC or RISC?

simple small instruction set



Shahrya

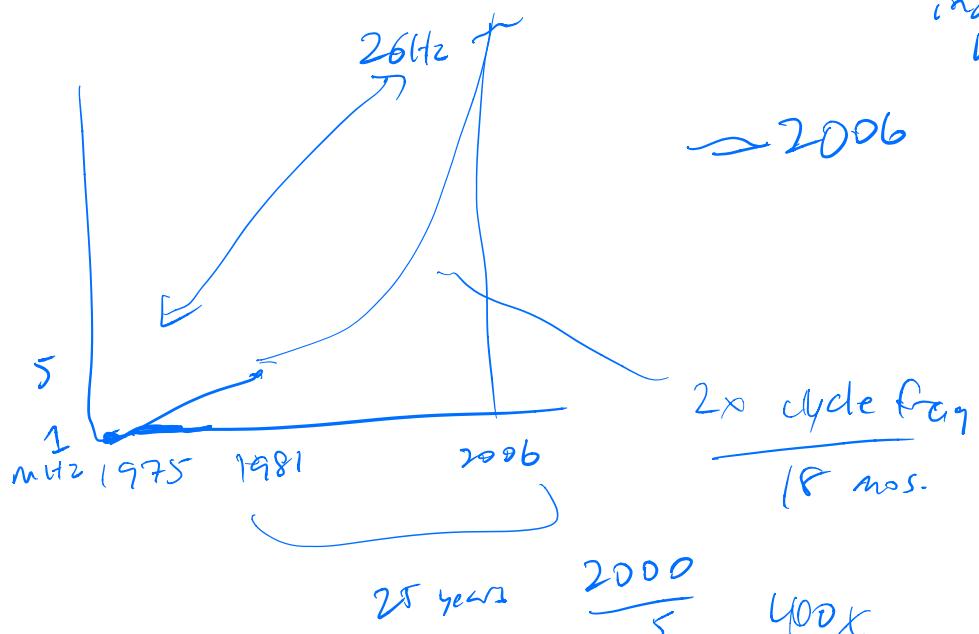


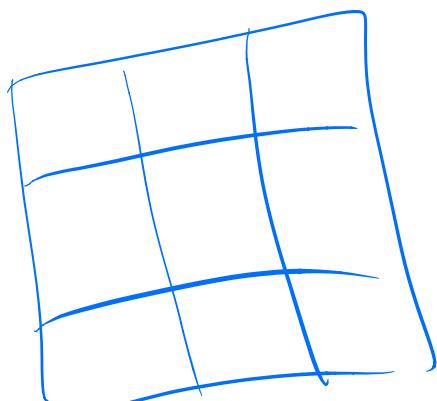
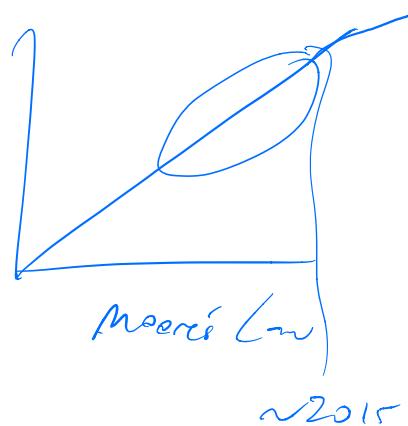
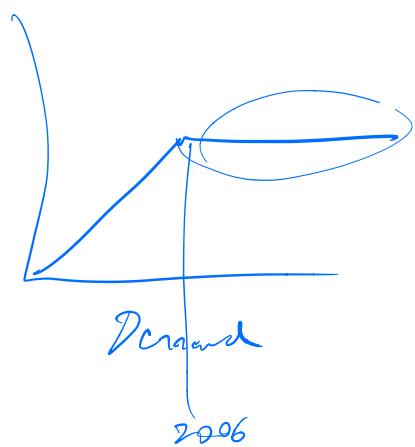
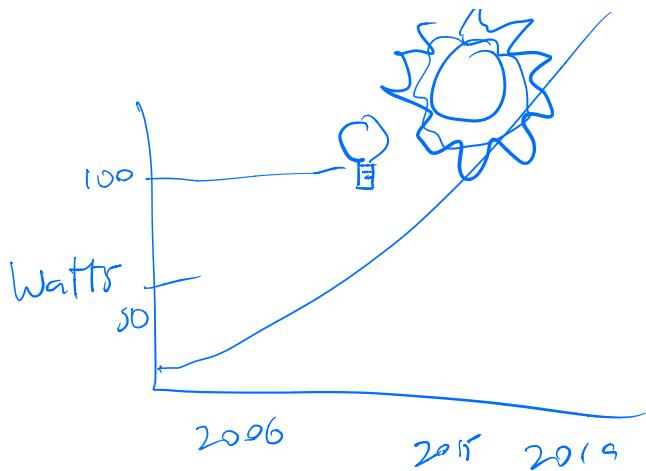
AMD

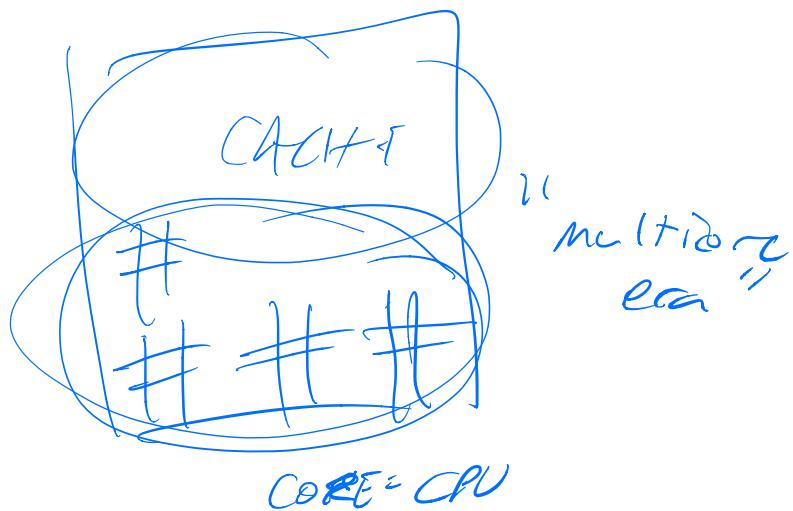
- fabs - not as good
- ⇒ low yield
- lower profits

→ 32-bit to 64-bit

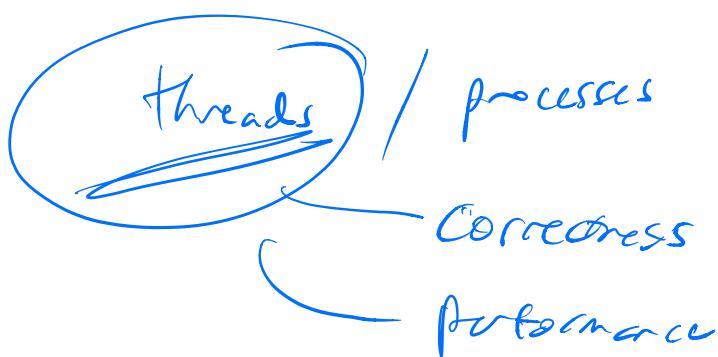
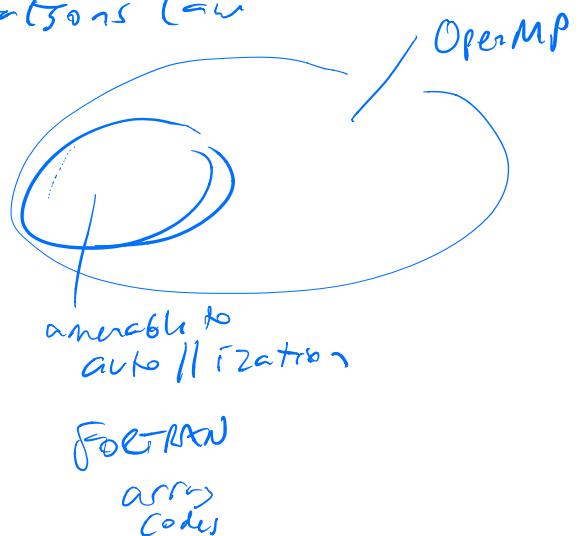
≈ multizone

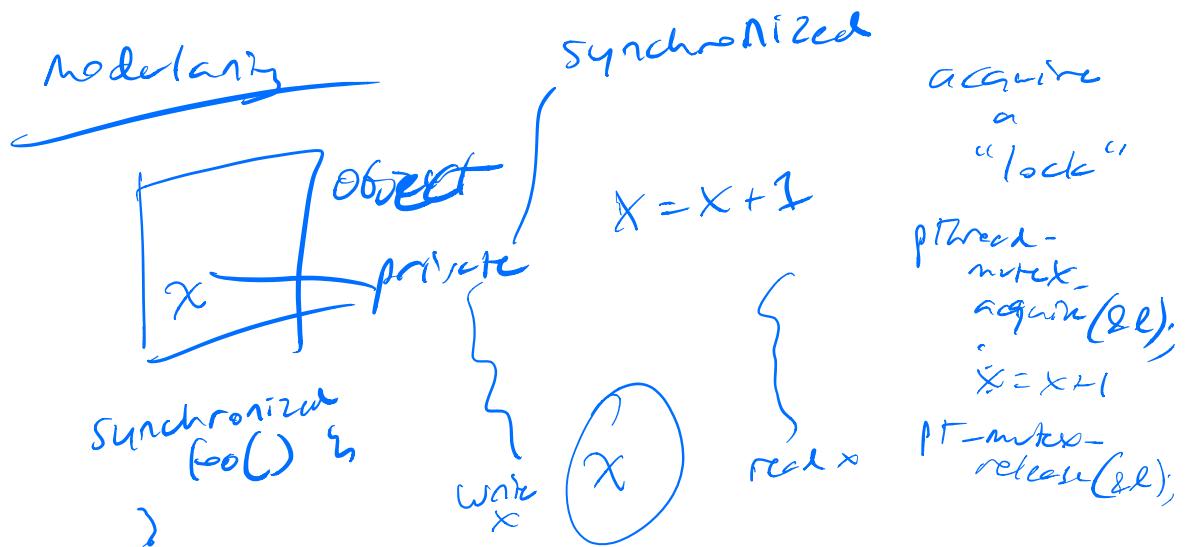






Amdahl's Law  
Gustavson's Law





"atomicity"

transactions: all or nothing

atoms - indivisible

atomic

Demostrates

transaction

- commit
- abort
- rollsback

"conflict"

retry

spin lock

```

lock
while (l == 1) // while locked
    ;
    ; // wait — sleep
(atomic) if (l == 0) { l = 1; } // acquire lock
TST & Lock — atomically

```

(unlock)  $l=0;$

"thundering herd"

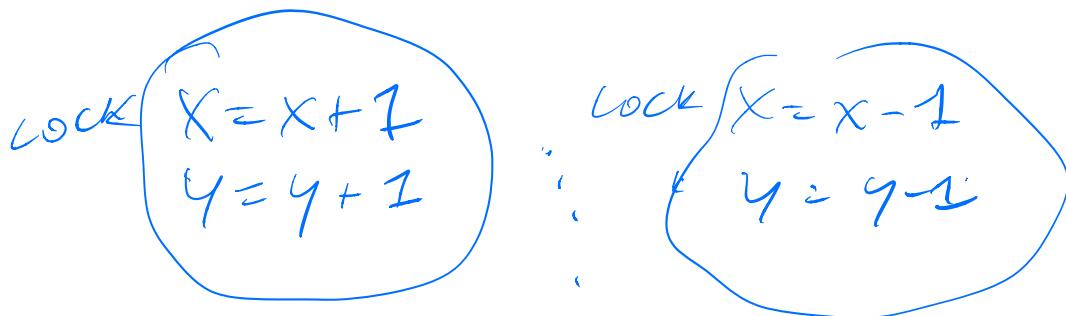
fair locks  
queuing locks

---

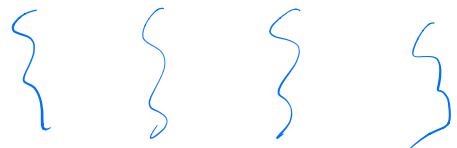
adding locks everywhere

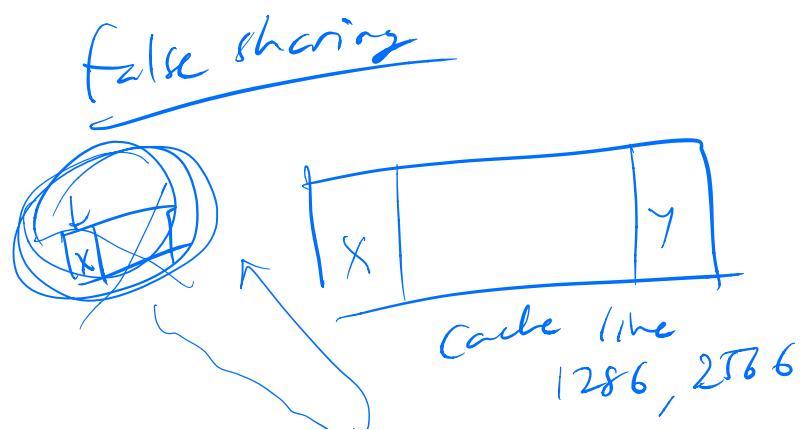
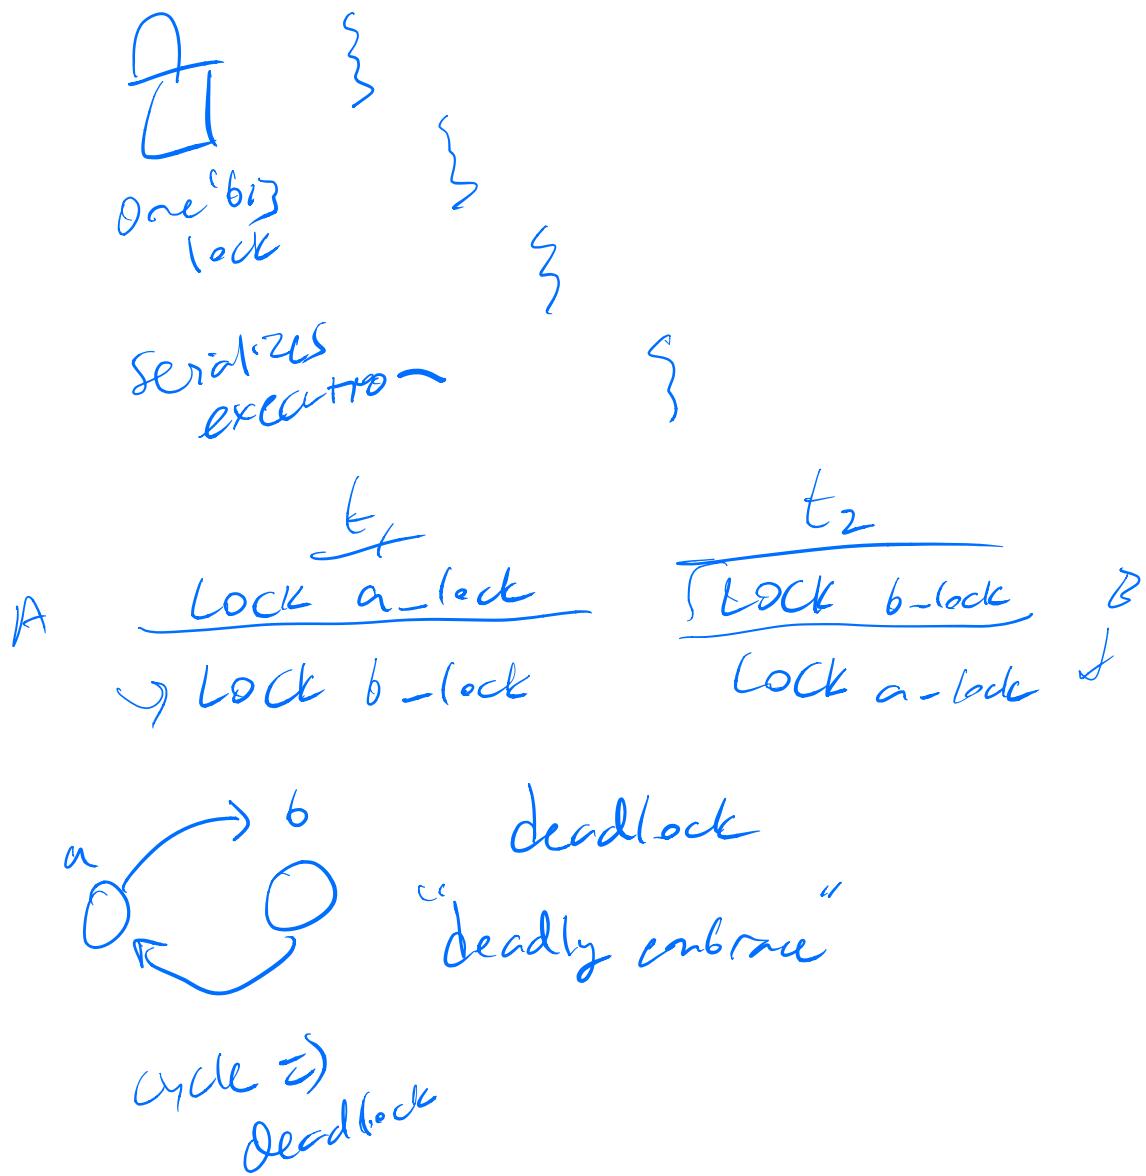
neither necessary (for correctness)  
nor sufficient (for performance)

$$x=0, y=0 \quad \left\{ \begin{array}{l} x=y \\ \end{array} \right.$$



fine-grained locks  $\rightarrow$  coarser  
for atomicity



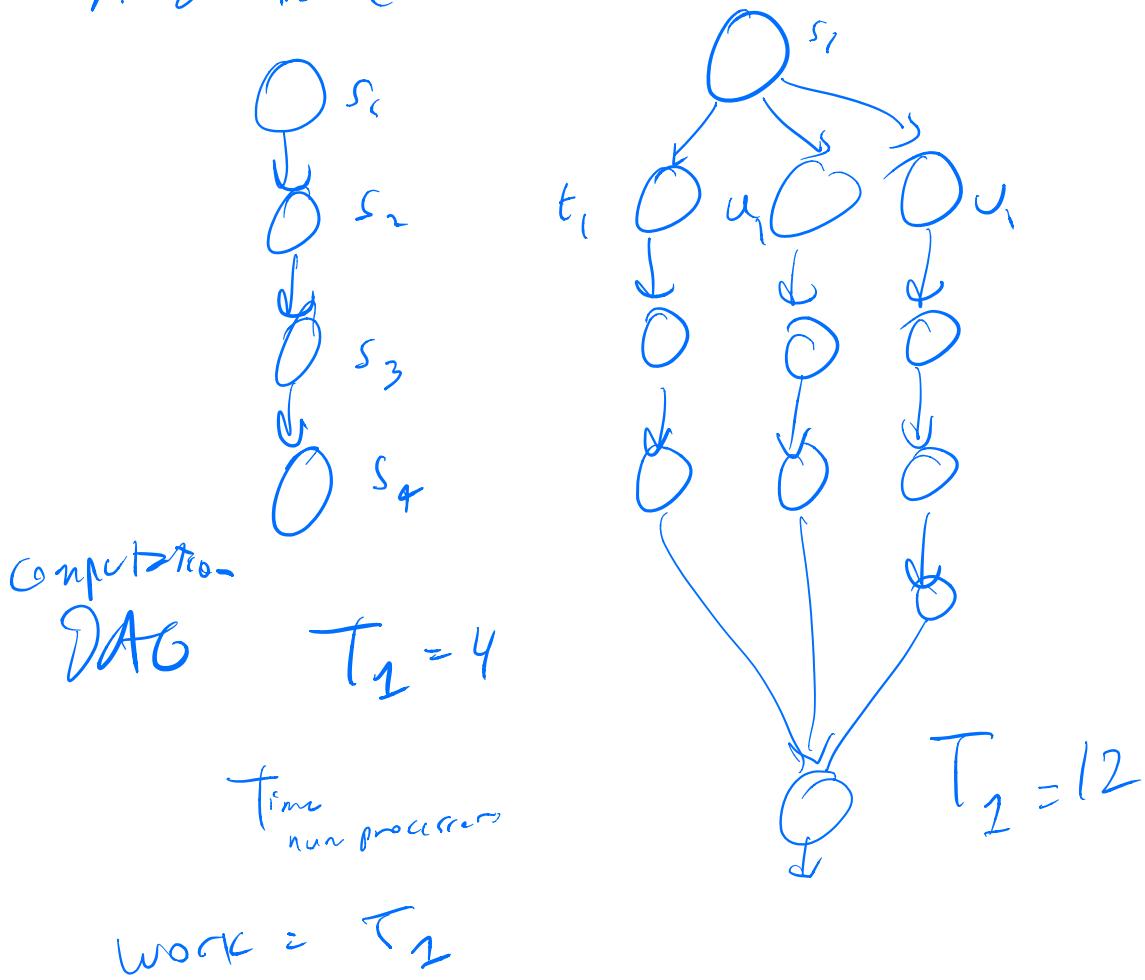


$\sim$

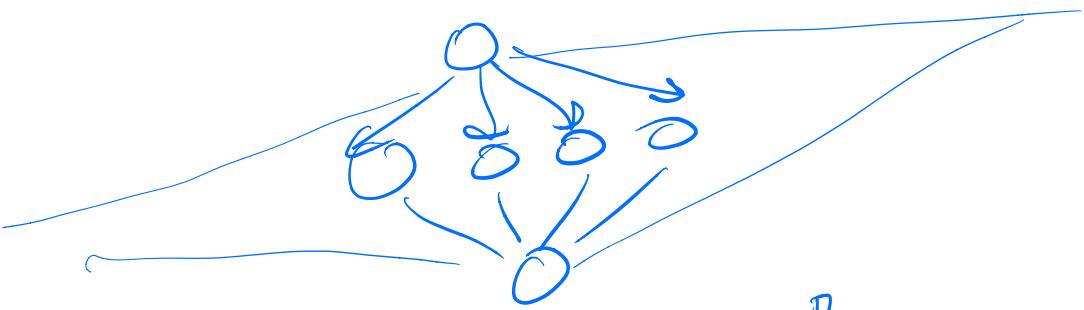
Cache line  
pins parsing

$\approx f\text{-}Ox$  slowdowns

Amdahl's Law



" $T_n \approx \frac{T_1}{n}$ " scale - perfect scaling



$T_{\text{co}}$  = control path  
 max from start to 6 from  
 "head of Data"

"serial part"

$$T_p \approx T_1 + T_{\text{co}}$$

$T_1$

parallel work

serial part

P

parallel work

90% of program — perfectly parallelizable

embarrassingly parallel

$$T_{\text{co}} = \frac{90\%}{P} + 10\%$$

0

100s → 50s

max speedup  $\rightarrow$  DX

VM hosting  $\approx$  embarrassingly II  
scales w/ # of instances

Amdahl's Law  
→ 1 workload fixed size

today's world  
Netflix → indep. customer moves  
emb. II  
 $\rightarrow$  # customers

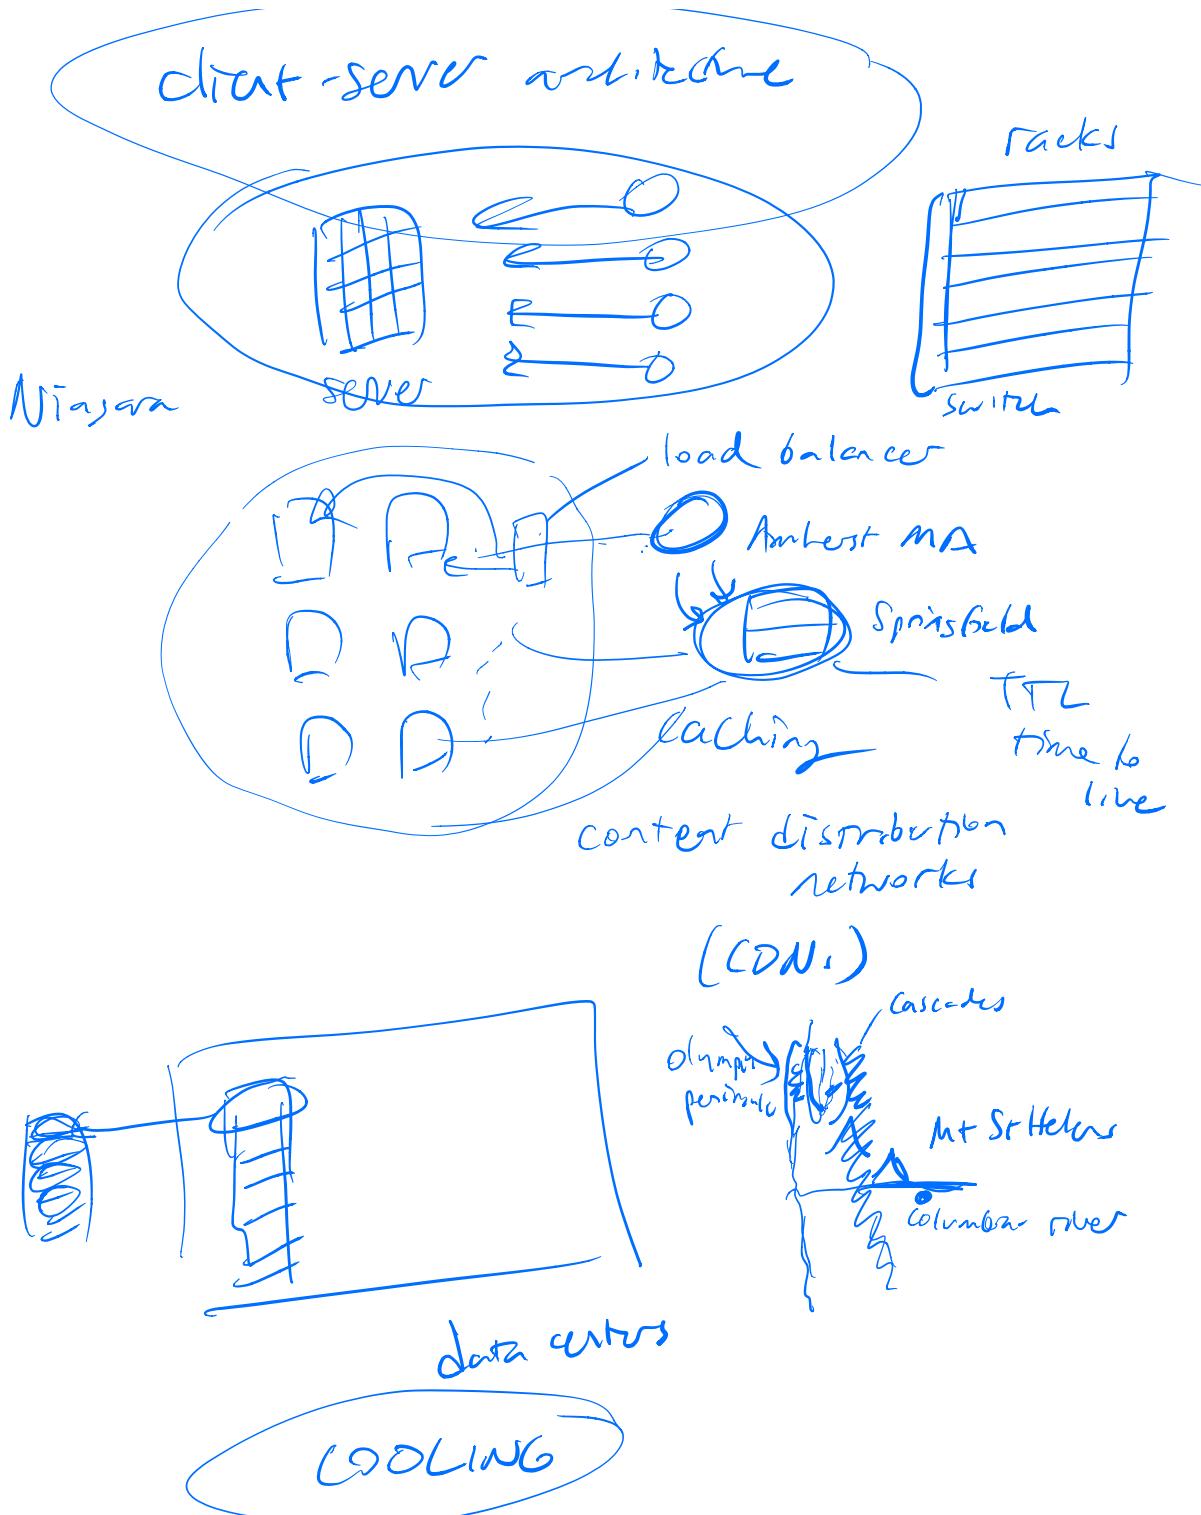
Google  
FB

Gustavson's Law

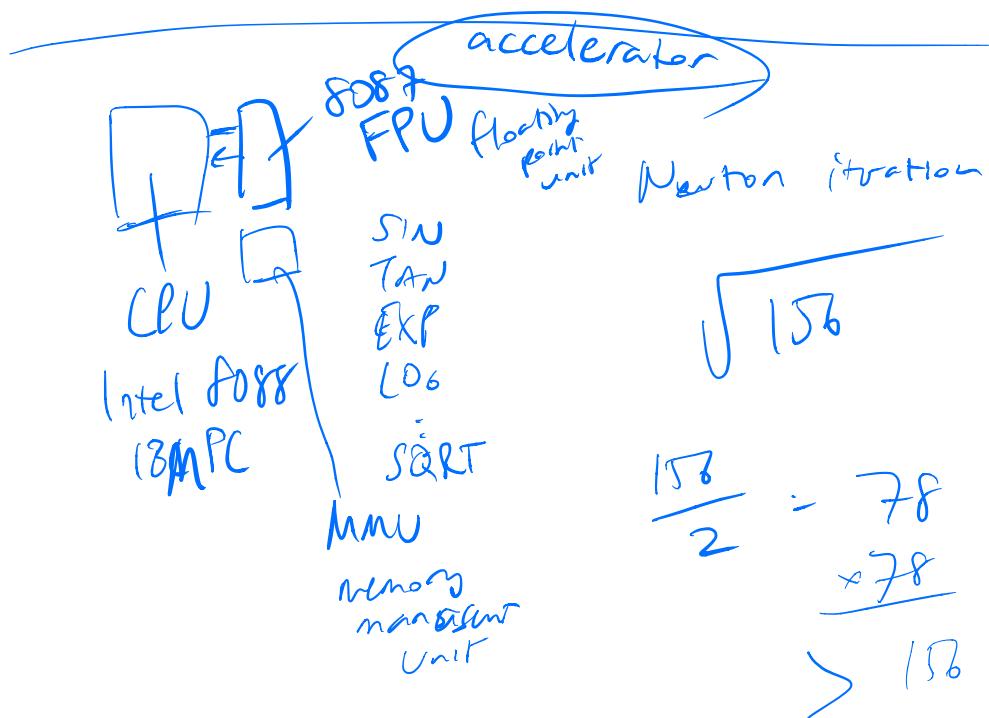
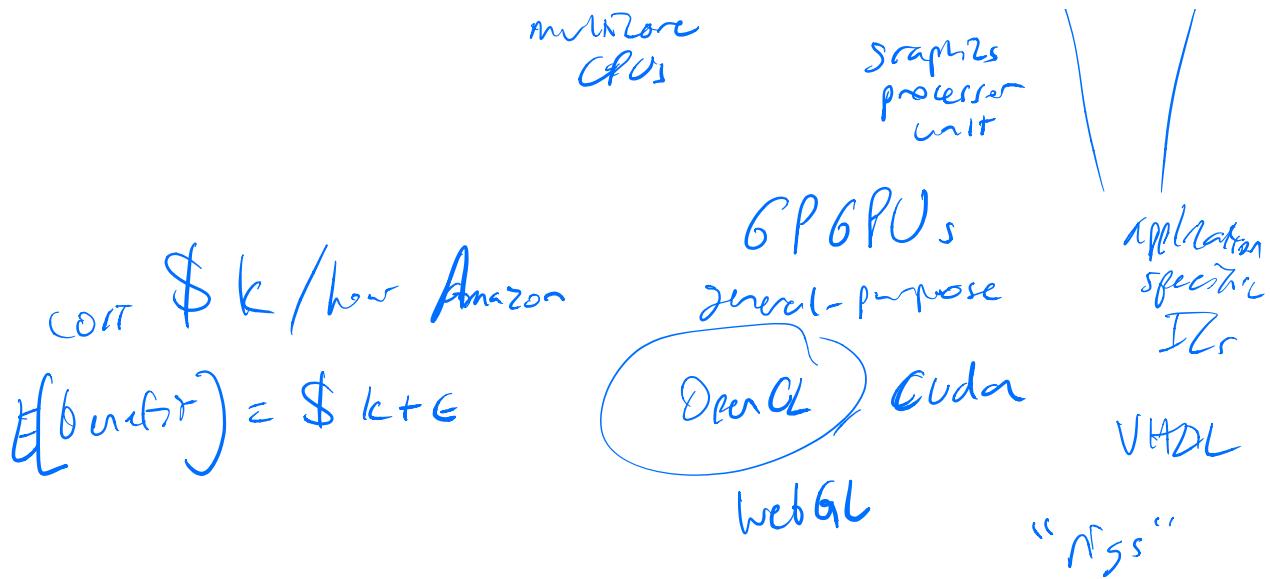
$$\Delta S \gg A_P$$

↑  
size of problem      ↑  
# of processors

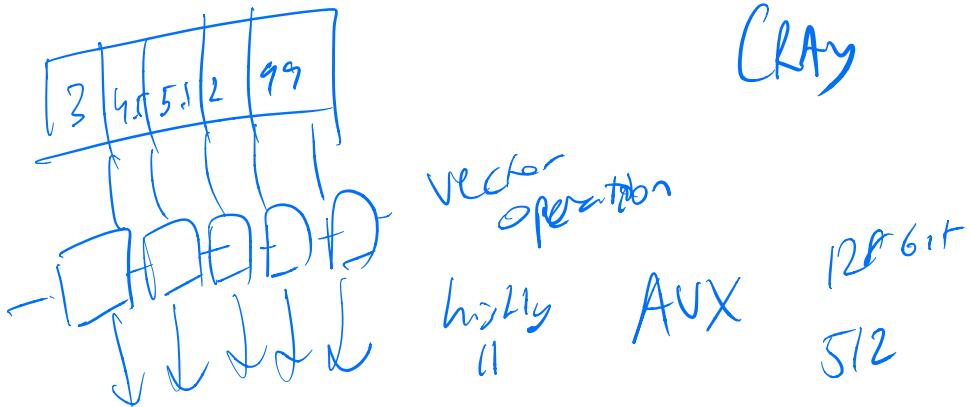
---



CPUs → cluster → GPUs → ASICs



Accelerators = CRYPTO  
Vectors



"Neural" ← HYPE

s/neural/matrix multiply/g

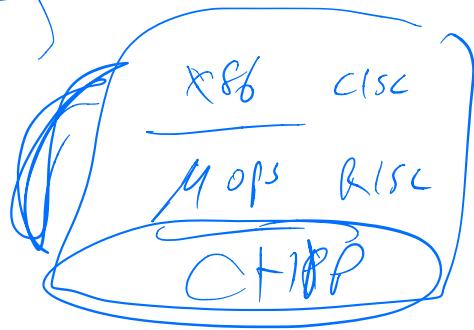
"tensor"

NPU  
TPU → Google  
tensor processing unit

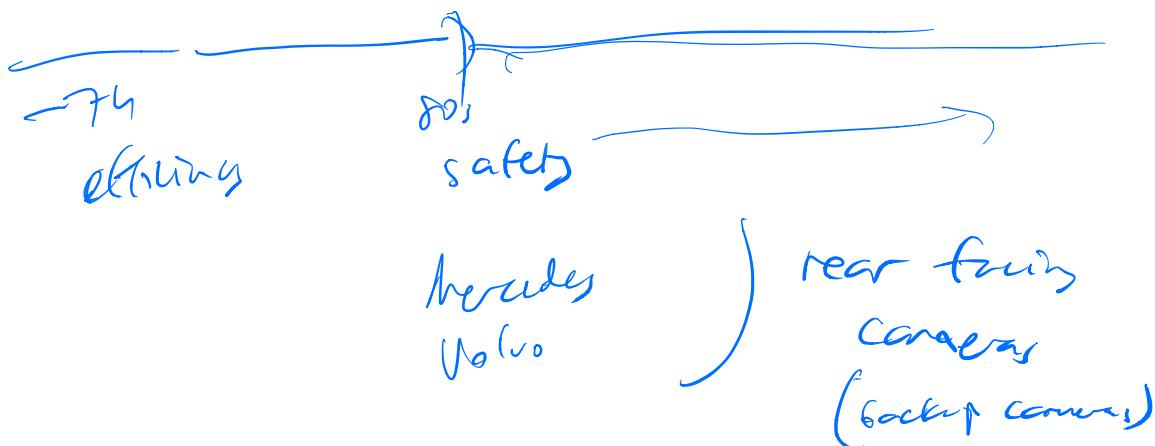
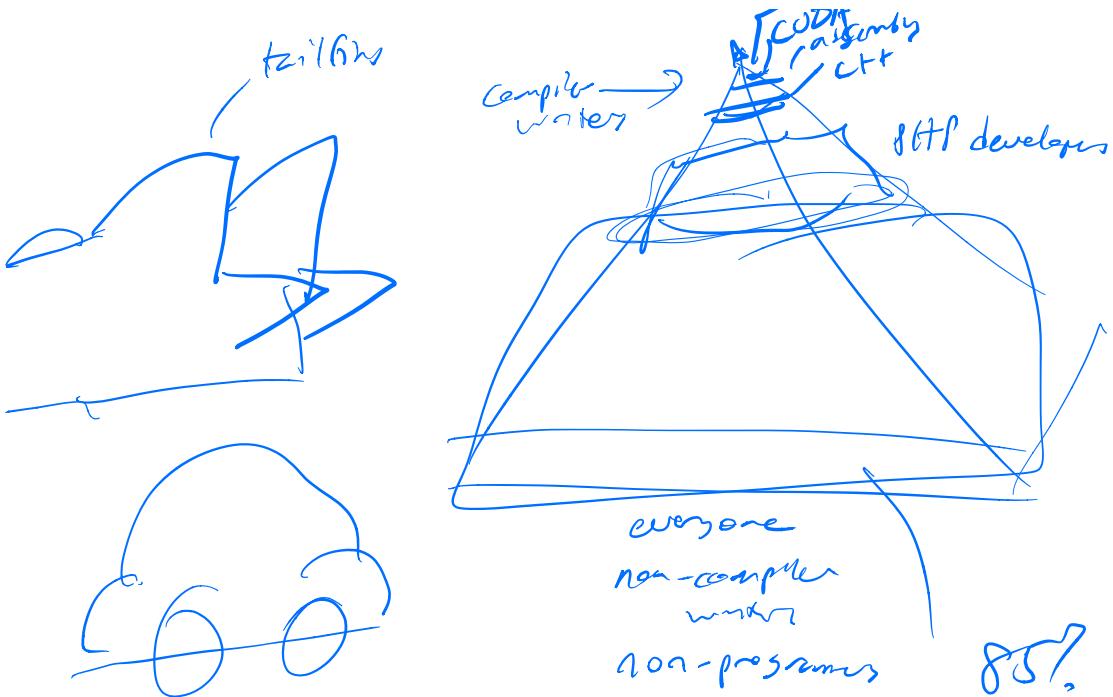
TVM

RISC vs. CISC

every Android Device  
→ ARM



four corelangs



~~17T H4  $\Rightarrow$  36T H1 - 3450 H8~~

~~Demand Scaling~~

phones

~~2 core~~  
~~4 core~~  
~~8 core~~

cameras  
potassium ≈  
energy / battery life

standard size instructions

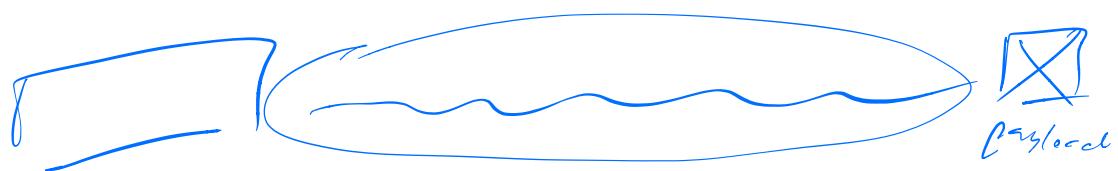
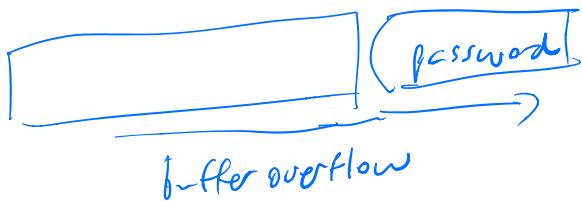
NOP → lock ROP stack — — —  
1 byte



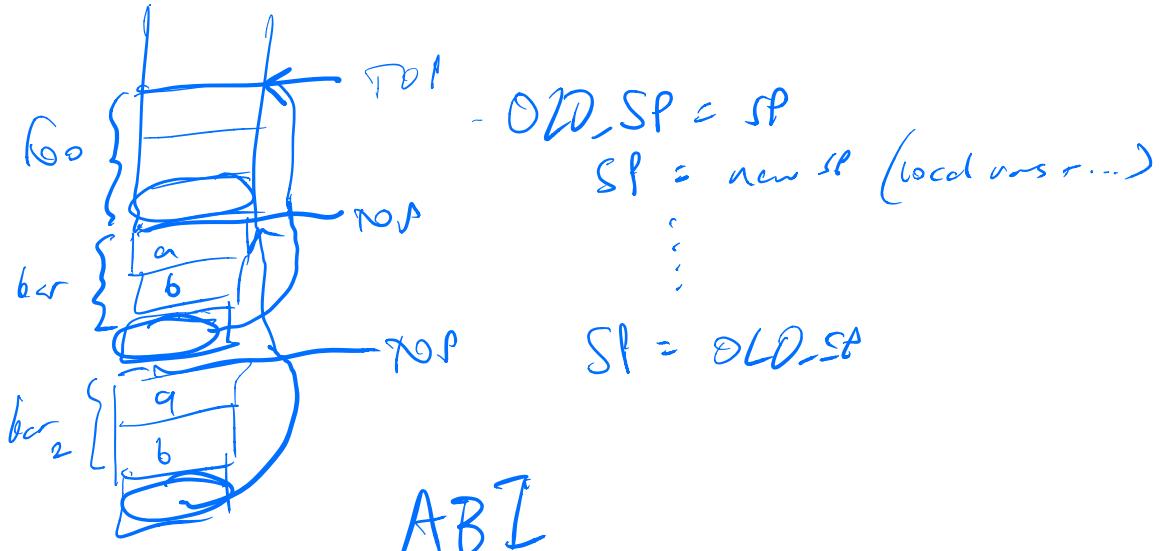
Variable sized instructions

ROP

low level security exploits — memory unsafety



## stack smashing attack



ABI

application binary interface



Morris worm 1988

DDoS distributed  
denial-of-service  
attack

NOP NOP NOP NOP , shellcode

The diagram illustrates a conceptual relationship. At the top, a blue double-headed arrow connects two concepts: 'Not sled' on the left and 'polymorphia' on the right. Below this, the word 'script kiddies' is written in blue, with a blue arrow pointing from it towards the 'Not sled' concept.

Dangling ptr / use-after-free

$x = \text{new } \text{fellow}$   
 $y = \dots \star$        $\boxed{\begin{matrix} x = y \\ \text{new} \end{matrix}} \star^?$   
 $;$   
 delete  $x$   
 $z = \text{new } -$   
 $y \rightarrow - -$       Use after  
 free by

"temporal safety"

class Foo {

    ~Foo() { }     } — destructor

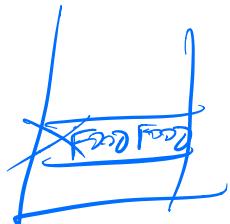
} ;

    } }  
    vtbl.  
    VMed  
    dispM  
    table

delete

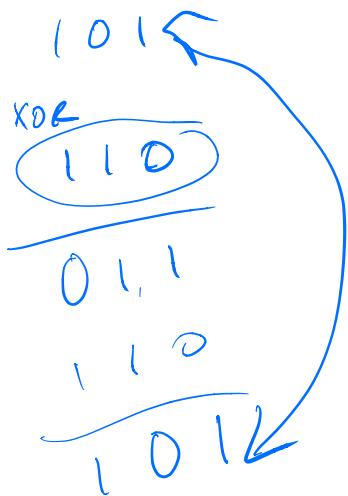
~~delete f;~~

stack smashing: shadow stack  
"Stack canaries"  
XOR ~~"key"~~



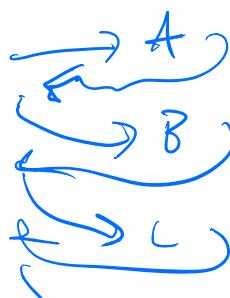
$x \wedge y$

$$\begin{array}{rcl} 1 & \text{xor } \phi & = 0 \\ 0 & \text{xor } 1 & = 1 \\ 1 & \text{xor } 0 & = 1 \\ 0 & \text{xor } 0 & = 0 \end{array}$$

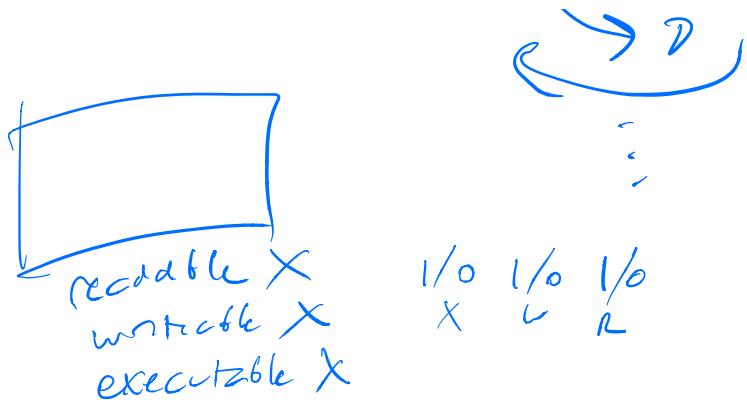


a xor key xor key =  $\alpha$

Dr Harder



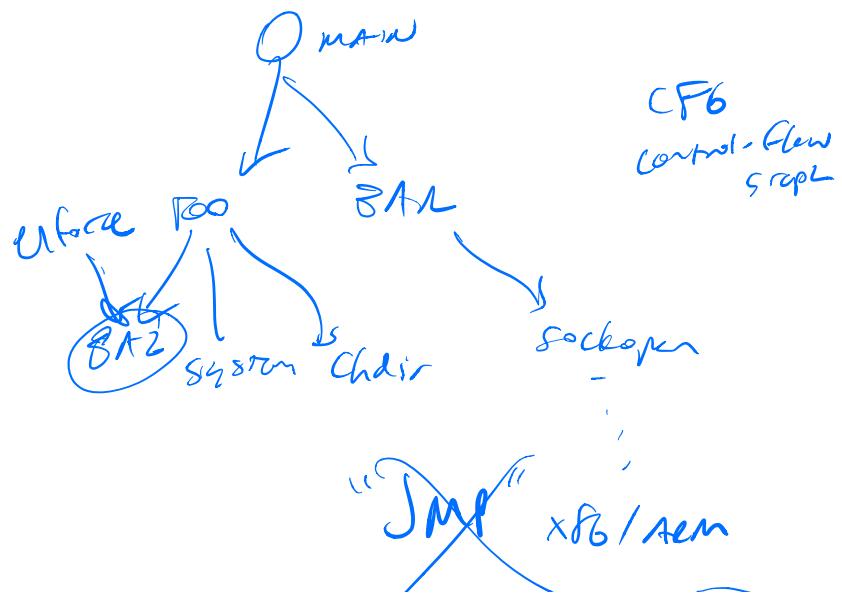
"gadgets"



## NX Protection

return-to-libc attack

"system"  
"rm -rf /"  
CFI      control-flow integrity

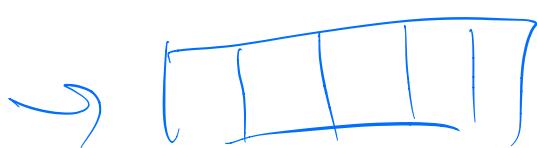




“One level of indirection”  
“all problems in CS”

RISC smaller instruction set  
“fixed-size instructions”

- ① easier to parse “decoder”
- ② harder to exploit
- ③ faster → easier to design how to make fast



PPC  
“TSX”  
HTM

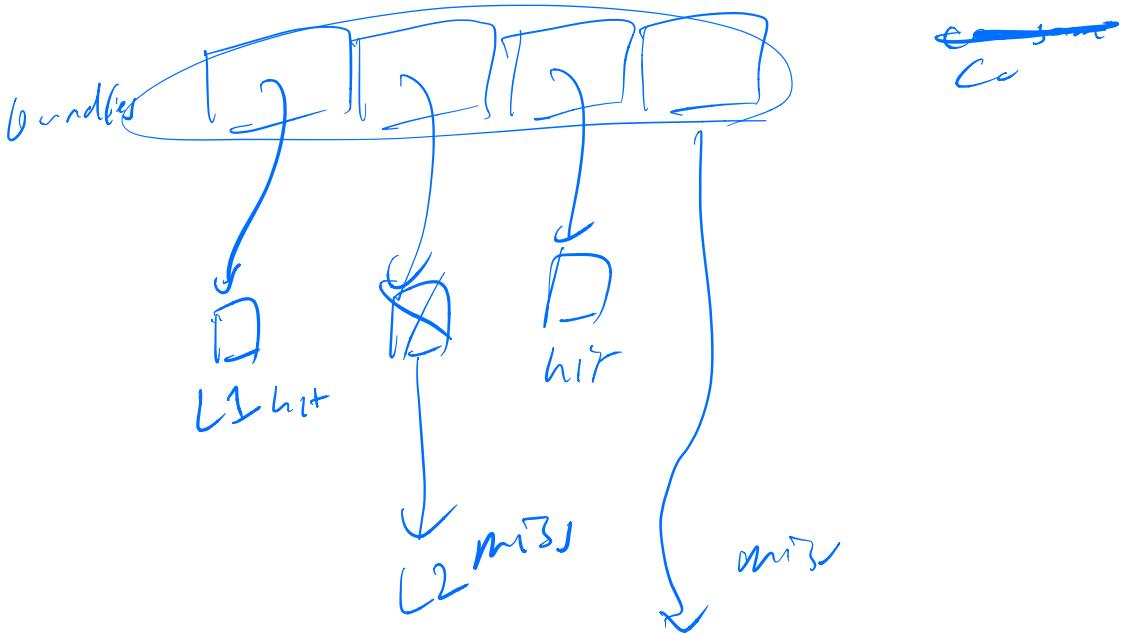
Amd - vector-instruction      Intel MPX  
- crypto      - memory protection

Google "ASan"  
AddressSanitizer

clang -fsanitize=address

VLIW

version instruction wed



EPIC

explicitly  
parallel  
instructions  
Co

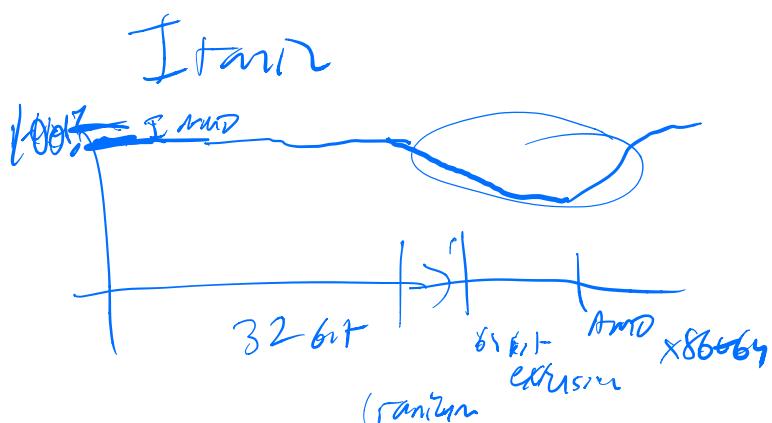
PREDICATION

(IF branch<sub>1</sub>)

DDDD

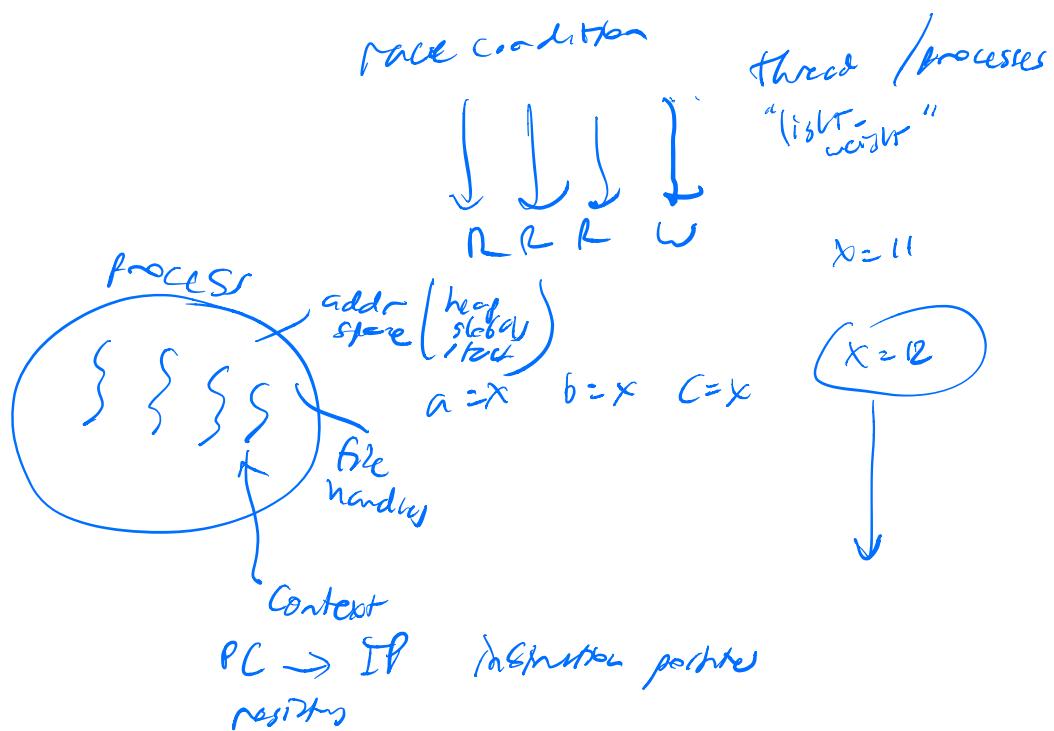
Itanium

IA-64



MOV  $\ddot{Q}$   
ADD  $\ddot{Q}$

## Concurrence



## Concurrence

~~errors~~  
races

atomicity violation

Order violation

deadlock

livelock

) condition variables

— wake-up problem

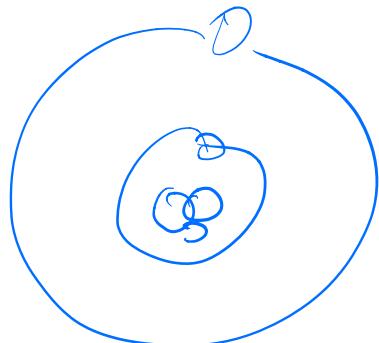
while(<sup>true</sup>  
    lock  
     $\rightarrow$  want()  
     $\wedge$  check state  
     $\wedge$   $T\theta(\cdot)$  break  
    }  
    })

NONDETERMINISTIC

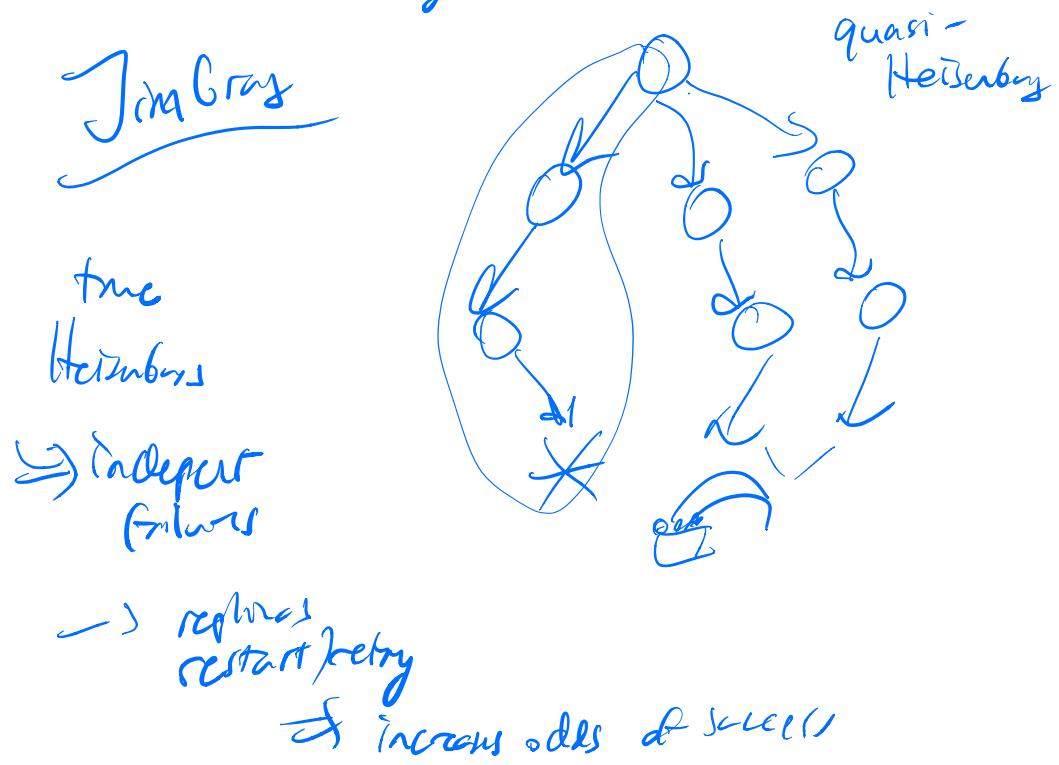
livelock

Heisenboss

DETERMINISTIC boss — Bohr boss



- severity?
- frequency?



Developers

debugging  
determinism!

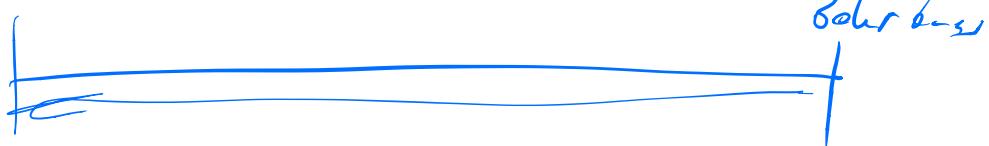
Bohr Bus

Users

Heisenberg

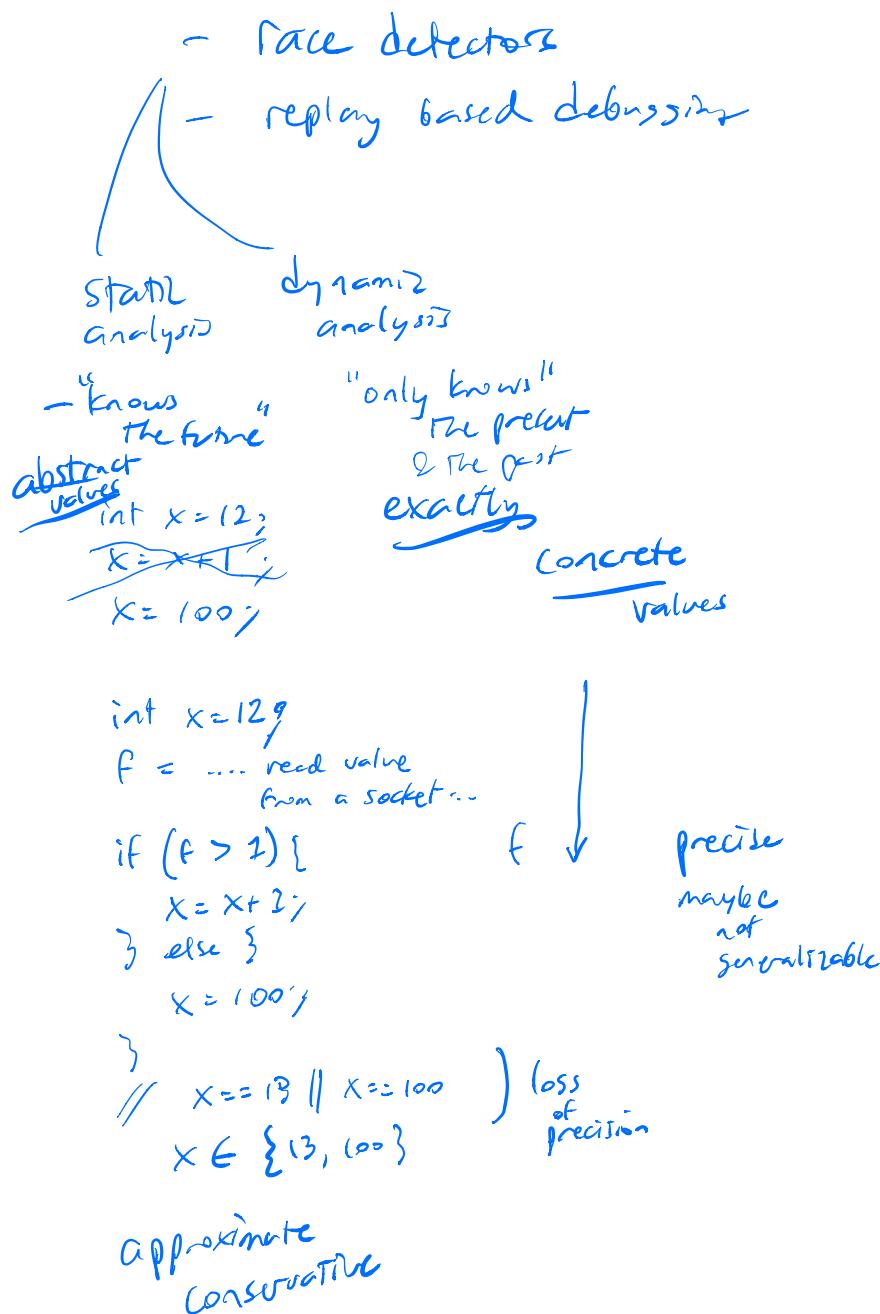
→ rely

→ fault tolerance  
(→ indep. failures)



D+E

100%



## bus detection

partial specification

P never derefs NULL

"implicit specification"

segfault buffer overflow  
NPE use-after-free  
races assertion failure

benign races / malignant races

"Hauswald!"

## formal verification

proves program P

implements spec S

code

y

log2

"complete / total  
specification"

abnormally



Word-sized  
read/write

X = 1

001

X = 4

100

101



double  
d

d = 0.1      d = 2.0

d = ...

" O-O

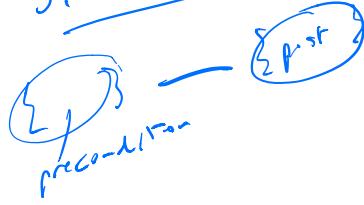
( ... )  
sort :  
postcondition      formal specifications

$H_i : 0 \leq i \leq N-1 :$   
 $a[i] \leq a[i+1]$

↓ " a is sorted  
in ascending order "

Postsort

Sir Tony Hoare



Quicksort  
average case  $O(n \log n)$   
worst case  $O(n^2)$

Mergesort

avg & worst case  $O(n \log n)$

quis custodes custodit

trusted computing base (TCB)

spec + solver  
+ code src  $\Rightarrow$  selected code  
TCB

$\approx 1000$  LOC  
just : hundreds !

targets for formal verification }  
Narrow APIs  
well specified  
really important

Project Everest  
"https" TLS F\*

## Heartbleed

static analysis for concurrency errors

~ false positives

precision

recall

print "NO BUG HERE!"

print "BUG!"

dyn analysis — no false positives

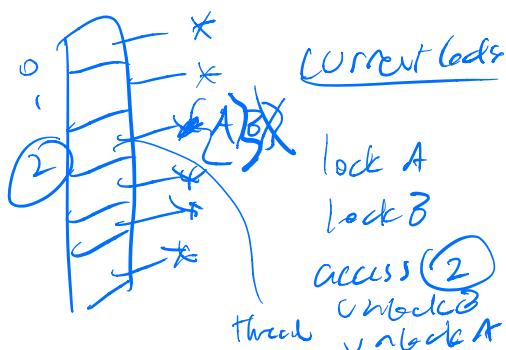
WITHACSS race!

— recall lower —

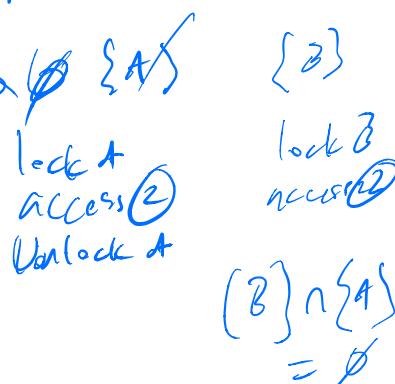
only reports race  
in that execution

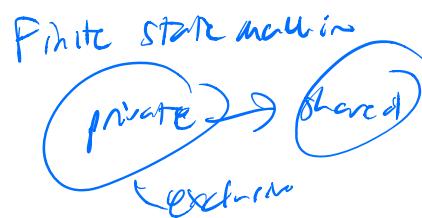
## Thread Sanitizer

### lockset analysis



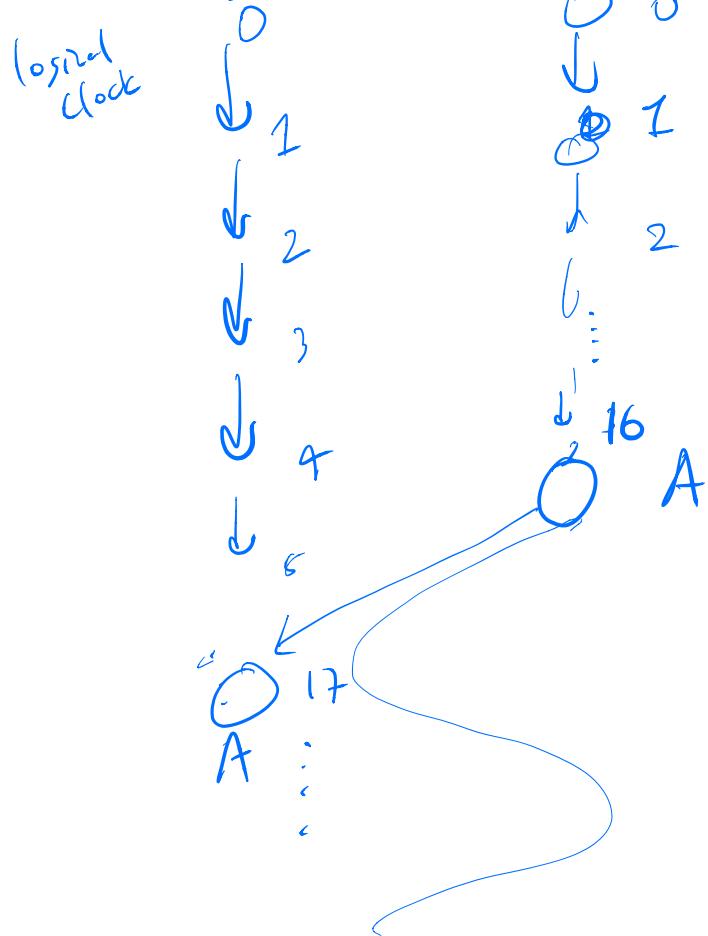
### happens-before analysis





Error

happens-before analysis



Why  
Threads?

- parallelism
  - coordination
  - hiding latency
    - I/O
    - memory
- ↑ performance

↳ average the periods  
↳ the delay until "first by"

"throughput" — rate of jobs safely done  
"throughput"

throughput-latency tradeoff

interactive

batch

|  
no context  
switches

CPU util ~ 100%

time  
sharing

/dev/fny

ed

MULTICS "Eunuchs"

1964 - 1967

MIT + GE + Bell Labs

- hierarchical FS

all previous:  
"flat"

relative  
paths



